

arbash-malik-assignment-2

April 5, 2024

0.1 Loading helper libraries

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import tensorflow as tf
import keras

from sklearn.model_selection import train_test_split

from keras.models import Sequential, Model
from keras.layers import BatchNormalization, Conv2D, Dense, Dropout, Flatten, GlobalAveragePooling2D, MaxPooling2D, Rescaling
from keras.callbacks import EarlyStopping
from keras.optimizers import Adam
from keras.applications import MobileNetV2
from keras.losses import MeanSquaredError
from keras.utils import to_categorical, set_random_seed

prng = np.random.RandomState(20240405)
set_random_seed(20240405)
keras.utils.set_random_seed(20240405)
```

0.2 Defining helper functions

```
[2]: def plot_model_history(model_histories, labels, main_title):
    plt.figure(figsize=(14, 6))

    # Adding main title
    plt.suptitle(main_title, fontsize=14)

    # plotting training & validation accuracy
    plt.subplot(1, 2, 1)
    for model_history, label in zip(model_histories, labels):
        epochs = range(1, len(model_history.history['accuracy']) + 1)
        plt.plot(epochs, model_history.history['accuracy'], label=f'{label} - Training')
```

```

        plt.plot(epochs, model_history.history['val_accuracy'], label=f'{label} - Validation', linestyle="--")
        plt.title('Accuracy')
        plt.xlabel('Epochs')
        plt.ylabel('Accuracy')
        plt.legend()

    # plotting training & validation loss
    plt.subplot(1, 2, 2)
    for model_history, label in zip(model_histories, labels):
        epochs = range(1, len(model_history.history['loss']) + 1)
        plt.plot(epochs, model_history.history['loss'], label=f'{label} - Training')
        plt.plot(epochs, model_history.history['val_loss'], label=f'{label} - Validation', linestyle="--")
        plt.title('Loss')
        plt.xlabel('Epochs')
        plt.ylabel('Loss')
        plt.legend()

    plt.tight_layout()
    plt.show()

```

0.3 Loading Data

```

[84]: from keras.datasets import fashion_mnist

(X_train, y_train), (X_test, y_test) = fashion_mnist.load_data()
X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, test_size=0.
    ↪2, random_state=prng)

print(f"X_train: {X_train.shape}")
print(f"y_train: {y_train.shape}")
print('__')
print(f"X_val: {X_val.shape}")
print(f"y_val: {y_val.shape}")
print('__')
print(f"X_test: {X_test.shape}")
print(f"y_test: {y_test.shape}")

```

X_train: (48000, 28, 28)

y_train: (48000,)

--

X_val: (12000, 28, 28)

y_val: (12000,)

--

X_test: (10000, 28, 28)

```
y_test: (10000,)
```

0.4 1. What would be an appropriate metric to evaluate your models? Why?

An appropriate metric to evaluate models for the classification task on this dataset would be **Accuracy**. Accuracy measures the proportion of correctly classified instances out of the total instances. Accuracy is also very easy to interpret. It represents the percentage of correct predictions made by the model. Accuracy is also relevant in our case as it directly measures how well the model is performing (the goal is to correctly classify images into one of the ten categories)

0.5 2. Get the data and show some example images from the data.

0.5.1 a) Training Data

```
[4]: class_names = ["T-shirt/  
    ↪top", "Trouser", "Pullover", "Dress", "Coat", "Sandal", "Shirt", "Sneaker", "Bag", "Ankle_  
    ↪boot"]  
  
fig, axs = plt.subplots(2, 5, figsize=(12,5))  
for i, ax in enumerate(axs.flatten()):  
    ax.imshow(X_train[i], cmap="binary")  
    ax.axis("off")  
    ax.set_title(f"Label: {class_names[y_train[i]]}")  
plt.tight_layout  
plt.show()
```



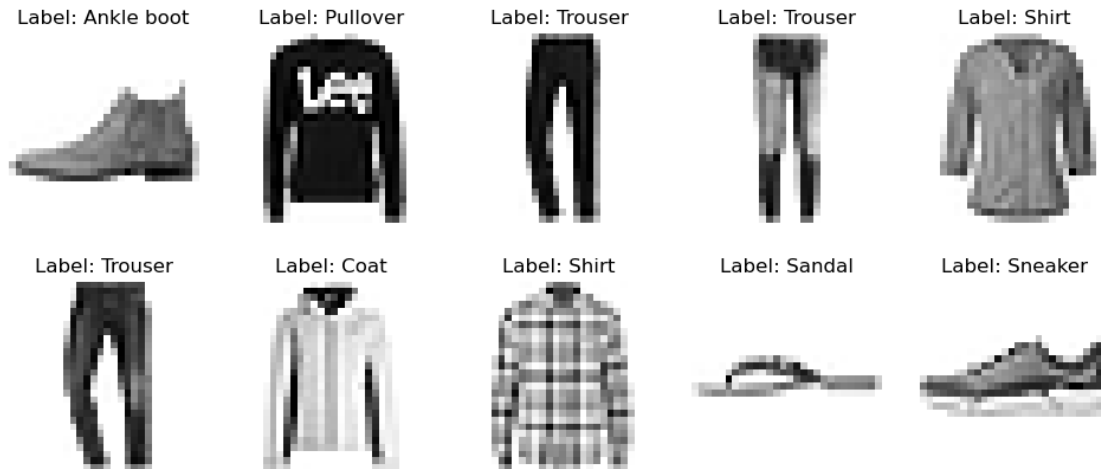
0.5.2 b) Validation Data

```
[5]: fig, axs = plt.subplots(2, 5, figsize=(12,5))
     for i, ax in enumerate(axs.flatten()):
         ax.imshow(X_val[i], cmap="binary")
         ax.axis("off")
         ax.set_title(f"Label: {class_names[y_val[i]]}")
     plt.tight_layout
     plt.show()
```



0.5.3 c) Test Data

```
[6]: fig, axs = plt.subplots(2, 5, figsize=(12,5))
     for i, ax in enumerate(axs.flatten()):
         ax.imshow(X_test[i], cmap="binary")
         ax.axis("off")
         ax.set_title(f"Label: {class_names[y_test[i]]}")
     plt.tight_layout
     plt.show()
```



0.6 3. Train a simple fully connected single hidden layer network to predict the items.

0.6.1 Remember to normalize the data similar to what we did in class. Make sure that you use enough epochs so that the validation error begins to level off - provide a plot of the training history.

```
[85]: from keras.utils import to_categorical

class_names = ["T-shirt/
↳top", "Trouser", "Pullover", "Dress", "Coat", "Sandal", "Shirt", "Sneaker", "Bag", "Ankle_
↳boot"]
num_classes = len(class_names)
print(f"Dimension of y before transformation: {y_train.shape}")

# Convert target variables to categorical
y_sets = [y_train, y_test, y_val]
y_train, y_test, y_val = [to_categorical(y, num_classes=num_classes) for y in
↳y_sets]
print(f"Dimension of y after transformation: {y_train.shape}")
```

Dimension of y before transformation: (48000,)
Dimension of y after transformation: (48000, 10)

```
[86]: from keras.models import Sequential
from keras.layers import Input, Flatten, Rescaling, Dense

model = Sequential([
    Input(shape=X_train.shape[1:]),
```

```

    Flatten(),
    Rescaling(1./255),
    Dense(100, activation='relu'),
    Dense(num_classes, activation='softmax')
])
print(model.summary())

# Compile the model
model.compile(loss='categorical_crossentropy', optimizer='adam',
              metrics=['accuracy'])

```

Model: "sequential_24"

Layer (type)	Output Shape	
Param #		
flatten_17 (Flatten)	(None, 784)	
↪ 0		
rescaling_18 (Rescaling)	(None, 784)	
↪ 0		
dense_42 (Dense)	(None, 100)	
↪ 78,500		
dense_43 (Dense)	(None, 10)	
↪ 1,010		

Total params: 79,510 (310.59 KB)

Trainable params: 79,510 (310.59 KB)

Non-trainable params: 0 (0.00 B)

None

```

[87]: history = model.fit(X_train, y_train, validation_data=(X_val, y_val),
                        epochs=20, batch_size = 512)

```

Epoch 1/20

94/94

1s 6ms/step -

accuracy: 0.6156 - loss: 1.1669 - val_accuracy: 0.8188 - val_loss: 0.5430


Epoch 2/20
94/94 0s 4ms/step -
accuracy: 0.8250 - loss: 0.5251 - val_accuracy: 0.8381 - val_loss: 0.4715
Epoch 3/20
94/94 0s 4ms/step -
accuracy: 0.8425 - loss: 0.4630 - val_accuracy: 0.8463 - val_loss: 0.4396
Epoch 4/20
94/94 0s 4ms/step -
accuracy: 0.8523 - loss: 0.4298 - val_accuracy: 0.8538 - val_loss: 0.4189
Epoch 5/20
94/94 0s 4ms/step -
accuracy: 0.8599 - loss: 0.4066 - val_accuracy: 0.8583 - val_loss: 0.4049
Epoch 6/20
94/94 0s 4ms/step -
accuracy: 0.8655 - loss: 0.3893 - val_accuracy: 0.8613 - val_loss: 0.3934
Epoch 7/20
94/94 0s 4ms/step -
accuracy: 0.8701 - loss: 0.3749 - val_accuracy: 0.8652 - val_loss: 0.3840
Epoch 8/20
94/94 0s 4ms/step -
accuracy: 0.8734 - loss: 0.3631 - val_accuracy: 0.8672 - val_loss: 0.3759
Epoch 9/20
94/94 0s 4ms/step -
accuracy: 0.8769 - loss: 0.3527 - val_accuracy: 0.8704 - val_loss: 0.3694
Epoch 10/20
94/94 0s 4ms/step -
accuracy: 0.8804 - loss: 0.3433 - val_accuracy: 0.8728 - val_loss: 0.3638
Epoch 11/20
94/94 0s 4ms/step -
accuracy: 0.8834 - loss: 0.3345 - val_accuracy: 0.8742 - val_loss: 0.3576
Epoch 12/20
94/94 0s 4ms/step -
accuracy: 0.8871 - loss: 0.3259 - val_accuracy: 0.8767 - val_loss: 0.3528
Epoch 13/20
94/94 0s 4ms/step -
accuracy: 0.8893 - loss: 0.3182 - val_accuracy: 0.8776 - val_loss: 0.3485
Epoch 14/20
94/94 0s 4ms/step -
accuracy: 0.8914 - loss: 0.3112 - val_accuracy: 0.8781 - val_loss: 0.3459
Epoch 15/20
94/94 0s 4ms/step -
accuracy: 0.8926 - loss: 0.3056 - val_accuracy: 0.8787 - val_loss: 0.3435
Epoch 16/20
94/94 0s 4ms/step -
accuracy: 0.8947 - loss: 0.3000 - val_accuracy: 0.8791 - val_loss: 0.3418
Epoch 17/20
94/94 0s 4ms/step -
accuracy: 0.8961 - loss: 0.2953 - val_accuracy: 0.8794 - val_loss: 0.3392

```
Epoch 18/20
94/94          0s 4ms/step -
accuracy: 0.8983 - loss: 0.2901 - val_accuracy: 0.8803 - val_loss: 0.3367
Epoch 19/20
94/94          0s 4ms/step -
accuracy: 0.8998 - loss: 0.2852 - val_accuracy: 0.8811 - val_loss: 0.3355
Epoch 20/20
94/94          0s 4ms/step -
accuracy: 0.9012 - loss: 0.2810 - val_accuracy: 0.8813 - val_loss: 0.3334
```

```
[88]: # Evaluation of the model on the validation set
scores = model.evaluate(X_val, y_val)
print(f"Accuracy for Basic Model: {round(scores[1], 4)}, Loss for Basic Model: {round(scores[0], 4)}")
```

```
375/375          1s 1ms/step -
accuracy: 0.8842 - loss: 0.3255
Accuracy for Basic Model: 0.8813, Loss for Basic Model: 0.3309
```

```
[89]: plot_model_history([history], ['Basic Model'], 'Basic Model')
```



0.7 4. Experiment with different network architectures and settings (number of hidden layers, number of nodes, regularization, etc.)

0.7.1 Train at least 3 models. Explain what you have tried and how it worked.

What I will iterate in my models are node numbers in each layer, adding hidden layers, adding dropout, adding early stopping clause.

Increasing the number of nodes allows for the model to capture the relationship between images and labels more clearly. Adding a hidden layer in a neural network increases the model's capacity

to learn complex patterns and representations from the input data.

Model 1: Nodes increased to 256 on first layer, added a second hidden layer with 100 nodes.

```
[90]: model1 = Sequential([
    Input(shape=X_train.shape[1:]),
    Flatten(),
    Rescaling(1./255),
    Dense(256, activation='relu'),
    Dense(100, activation='relu'),
    Dense(num_classes, activation='softmax')
])
print(model1.summary())

# Compile the model
model1.compile(loss='categorical_crossentropy', optimizer='adam',
               metrics=['accuracy'])
```

Model: "sequential_25"

Layer (type)	Output Shape	
Param #		
flatten_18 (Flatten)	(None, 784)	
↪ 0		
rescaling_19 (Rescaling)	(None, 784)	
↪ 0		
dense_44 (Dense)	(None, 256)	
↪200,960		
dense_45 (Dense)	(None, 100)	
↪25,700		
dense_46 (Dense)	(None, 10)	
↪1,010		

Total params: 227,670 (889.34 KB)

Trainable params: 227,670 (889.34 KB)

Non-trainable params: 0 (0.00 B)

None

```
[91]: history1 = model1.fit(X_train, y_train, validation_data=(X_val, y_val),  
    ↪ epochs=20, batch_size = 512)
```

```
Epoch 1/20  
94/94          2s 8ms/step -  
accuracy: 0.6732 - loss: 0.9736 - val_accuracy: 0.8355 - val_loss: 0.4664  
Epoch 2/20  
94/94          1s 6ms/step -  
accuracy: 0.8434 - loss: 0.4493 - val_accuracy: 0.8580 - val_loss: 0.4026  
Epoch 3/20  
94/94          1s 6ms/step -  
accuracy: 0.8604 - loss: 0.3944 - val_accuracy: 0.8702 - val_loss: 0.3727  
Epoch 4/20  
94/94          1s 6ms/step -  
accuracy: 0.8716 - loss: 0.3625 - val_accuracy: 0.8748 - val_loss: 0.3566  
Epoch 5/20  
94/94          1s 6ms/step -  
accuracy: 0.8807 - loss: 0.3392 - val_accuracy: 0.8793 - val_loss: 0.3441  
Epoch 6/20  
94/94          1s 6ms/step -  
accuracy: 0.8875 - loss: 0.3201 - val_accuracy: 0.8825 - val_loss: 0.3352  
Epoch 7/20  
94/94          1s 6ms/step -  
accuracy: 0.8911 - loss: 0.3049 - val_accuracy: 0.8810 - val_loss: 0.3286  
Epoch 8/20  
94/94          1s 6ms/step -  
accuracy: 0.8948 - loss: 0.2929 - val_accuracy: 0.8863 - val_loss: 0.3199  
Epoch 9/20  
94/94          1s 6ms/step -  
accuracy: 0.8990 - loss: 0.2823 - val_accuracy: 0.8900 - val_loss: 0.3114  
Epoch 10/20  
94/94          1s 6ms/step -  
accuracy: 0.9026 - loss: 0.2706 - val_accuracy: 0.8917 - val_loss: 0.3060  
Epoch 11/20  
94/94          1s 6ms/step -  
accuracy: 0.9046 - loss: 0.2621 - val_accuracy: 0.8924 - val_loss: 0.3023  
Epoch 12/20  
94/94          1s 6ms/step -  
accuracy: 0.9079 - loss: 0.2520 - val_accuracy: 0.8915 - val_loss: 0.3026  
Epoch 13/20  
94/94          1s 6ms/step -  
accuracy: 0.9111 - loss: 0.2454 - val_accuracy: 0.8900 - val_loss: 0.3062  
Epoch 14/20
```

```

94/94          1s 6ms/step -
accuracy: 0.9130 - loss: 0.2371 - val_accuracy: 0.8915 - val_loss: 0.3069
Epoch 15/20
94/94          1s 6ms/step -
accuracy: 0.9171 - loss: 0.2287 - val_accuracy: 0.8924 - val_loss: 0.3074
Epoch 16/20
94/94          1s 7ms/step -
accuracy: 0.9203 - loss: 0.2213 - val_accuracy: 0.8894 - val_loss: 0.3133
Epoch 17/20
94/94          1s 8ms/step -
accuracy: 0.9226 - loss: 0.2142 - val_accuracy: 0.8900 - val_loss: 0.3153
Epoch 18/20
94/94          1s 8ms/step -
accuracy: 0.9248 - loss: 0.2085 - val_accuracy: 0.8929 - val_loss: 0.3103
Epoch 19/20
94/94          1s 8ms/step -
accuracy: 0.9263 - loss: 0.2027 - val_accuracy: 0.8916 - val_loss: 0.3087
Epoch 20/20
94/94          1s 8ms/step -
accuracy: 0.9281 - loss: 0.1993 - val_accuracy: 0.8917 - val_loss: 0.3122

```

```

[92]: # Evaluation of the model on the validation set
scores1 = model1.evaluate(X_val, y_val)
print("\n")
print(f"Accuracy for Model 1: {round(scores1[1], 4)}, Loss for Model 1: ␣
      ↳{round(scores1[0], 4)}")

```

```

375/375          1s 2ms/step -
accuracy: 0.8960 - loss: 0.3036


```

Accuracy for Model 1: 0.8917, Loss for Model 1: 0.3103

```

[93]: plot_model_history([history, history1], ['Basic Model', 'Model 1'], 'Comparison -␣
      ↳Model vs Model 1')

```



Comparing to our basic model, our first iteration of the model where increased the nodes on the first layer and added a layer worked pretty well as both accuracy and loss improve

Model 2: 256 Nodes on first layer, added a second hidden layer with 100 nodes, added a third hidden layer with 50 nodes

```
[16]: model2 = Sequential([
    Input(shape=X_train.shape[1:]),
    Flatten(),
    Rescaling(1./255),
    Dense(256, activation='relu'),
    Dense(100, activation='relu'),
    Dense(50, activation='relu'),
    Dense(num_classes, activation='softmax')
])

# Compile the model
model2.compile(loss='categorical_crossentropy', optimizer='adam',
               metrics=['accuracy'])

print(model2.summary())
```

Model: "sequential_2"

Layer (type)

Output Shape

Param #

flatten_2 (Flatten)

(None, 784)

0

```

rescaling_2 (Rescaling)          (None, 784)
↳ 0

dense_5 (Dense)                  (None, 256)
↳200,960

dense_6 (Dense)                  (None, 100)
↳25,700

dense_7 (Dense)                  (None, 50)
↳5,050

dense_8 (Dense)                  (None, 10)
↳510

```

Total params: 232,220 (907.11 KB)

Trainable params: 232,220 (907.11 KB)

Non-trainable params: 0 (0.00 B)

None

```
[17]: history2 = model2.fit(X_train, y_train, validation_data=(X_val, y_val),
↳epochs=20, batch_size = 512)
```

```

Epoch 1/20
94/94          4s 17ms/step -
accuracy: 0.6416 - loss: 1.0570 - val_accuracy: 0.8225 - val_loss: 0.5074
Epoch 2/20
94/94          1s 12ms/step -
accuracy: 0.8382 - loss: 0.4625 - val_accuracy: 0.8390 - val_loss: 0.4593
Epoch 3/20
94/94          1s 11ms/step -
accuracy: 0.8543 - loss: 0.4182 - val_accuracy: 0.8641 - val_loss: 0.3879
Epoch 4/20
94/94          1s 13ms/step -
accuracy: 0.8642 - loss: 0.3822 - val_accuracy: 0.8685 - val_loss: 0.3720
Epoch 5/20
94/94          1s 13ms/step -
accuracy: 0.8734 - loss: 0.3568 - val_accuracy: 0.8701 - val_loss: 0.3631
Epoch 6/20
94/94          1s 14ms/step -
accuracy: 0.8792 - loss: 0.3360 - val_accuracy: 0.8734 - val_loss: 0.3522

```

```

Epoch 7/20
94/94          1s 12ms/step -
accuracy: 0.8848 - loss: 0.3201 - val_accuracy: 0.8753 - val_loss: 0.3422
Epoch 8/20
94/94          1s 13ms/step -
accuracy: 0.8884 - loss: 0.3074 - val_accuracy: 0.8798 - val_loss: 0.3374
Epoch 9/20
94/94          1s 13ms/step -
accuracy: 0.8925 - loss: 0.2943 - val_accuracy: 0.8783 - val_loss: 0.3371
Epoch 10/20
94/94          1s 12ms/step -
accuracy: 0.8965 - loss: 0.2839 - val_accuracy: 0.8827 - val_loss: 0.3239
Epoch 11/20
94/94          1s 12ms/step -
accuracy: 0.9012 - loss: 0.2710 - val_accuracy: 0.8862 - val_loss: 0.3197
Epoch 12/20
94/94          1s 13ms/step -
accuracy: 0.9025 - loss: 0.2639 - val_accuracy: 0.8843 - val_loss: 0.3231
Epoch 13/20
94/94          1s 13ms/step -
accuracy: 0.9062 - loss: 0.2560 - val_accuracy: 0.8817 - val_loss: 0.3342
Epoch 14/20
94/94          1s 15ms/step -
accuracy: 0.9080 - loss: 0.2488 - val_accuracy: 0.8773 - val_loss: 0.3479
Epoch 15/20
94/94          1s 13ms/step -
accuracy: 0.9114 - loss: 0.2404 - val_accuracy: 0.8773 - val_loss: 0.3472
Epoch 16/20
94/94          1s 13ms/step -
accuracy: 0.9152 - loss: 0.2298 - val_accuracy: 0.8827 - val_loss: 0.3312
Epoch 17/20
94/94          1s 12ms/step -
accuracy: 0.9187 - loss: 0.2200 - val_accuracy: 0.8834 - val_loss: 0.3334
Epoch 18/20
94/94          1s 13ms/step -
accuracy: 0.9220 - loss: 0.2139 - val_accuracy: 0.8845 - val_loss: 0.3295
Epoch 19/20
94/94          1s 12ms/step -
accuracy: 0.9254 - loss: 0.2083 - val_accuracy: 0.8821 - val_loss: 0.3341
Epoch 20/20
94/94          1s 13ms/step -
accuracy: 0.9247 - loss: 0.2060 - val_accuracy: 0.8796 - val_loss: 0.3519

```

```

[18]: # Evaluation of the model on the validation set
      scores2 = model2.evaluate(X_val, y_val)

```


```
print(f"Accuracy for Model 2: {round(scores2[1], 4)}, Loss for Model 2: {round(scores2[0], 4)}")

print("\n")
```

375/375 1s 3ms/step -
accuracy: 0.8781 - loss: 0.3501

Accuracy for Model 2: 0.8796, Loss for Model 2: 0.3535

```
[94]: plot_model_history([history1, history2], ['Model 1', 'Model 2'], 'Comparison - Model 1 vs Model 2')
```



Comparing to our first iteration model, our second iteration of the model where we added a third hidden layer is the worse in accuracy as well as the loss is worse in validation set

Model 3: 100 Nodes on first layer, added a second hidden layer with 25 nodes, added a third hidden layer with 50 nodes

```
[19]: model3 = Sequential([
    Input(shape=X_train.shape[1:]),
    Flatten(),
    Rescaling(1./255),
    Dense(100, activation='relu'),
    Dense(75, activation='relu'),
    Dense(50, activation='relu'),
    Dense(num_classes, activation='softmax')
])
print(model3.summary())
```

```
# Compile the model
model3.compile(loss='categorical_crossentropy', optimizer='adam',
               metrics=['accuracy'])
```

Model: "sequential_3"

Layer (type)	Output Shape	
Param #		
flatten_3 (Flatten)	(None, 784)	
↪ 0		
rescaling_3 (Rescaling)	(None, 784)	
↪ 0		
dense_9 (Dense)	(None, 100)	
↪ 78,500		
dense_10 (Dense)	(None, 75)	
↪ 7,575		
dense_11 (Dense)	(None, 50)	
↪ 3,800		
dense_12 (Dense)	(None, 10)	
↪ 510		

Total params: 90,385 (353.07 KB)

Trainable params: 90,385 (353.07 KB)

Non-trainable params: 0 (0.00 B)

None

```
[20]: history3 = model3.fit(X_train, y_train, validation_data=(X_val, y_val),
                           epochs=100, batch_size = 512
                           ,callbacks=[EarlyStopping(monitor='val_accuracy',
                           patience=10)])
```

Epoch 1/100

94/94 4s 16ms/step -
accuracy: 0.6037 - loss: 1.2183 - val_accuracy: 0.8240 - val_loss: 0.4995
Epoch 2/100

94/94 2s 17ms/step -
accuracy: 0.8270 - loss: 0.4932 - val_accuracy: 0.8407 - val_loss: 0.4456
Epoch 3/100

94/94 2s 20ms/step -
accuracy: 0.8488 - loss: 0.4281 - val_accuracy: 0.8549 - val_loss: 0.4106
Epoch 4/100

94/94 2s 16ms/step -
accuracy: 0.8606 - loss: 0.3963 - val_accuracy: 0.8641 - val_loss: 0.3821
Epoch 5/100

94/94 2s 23ms/step -
accuracy: 0.8697 - loss: 0.3699 - val_accuracy: 0.8708 - val_loss: 0.3646
Epoch 6/100

94/94 2s 16ms/step -
accuracy: 0.8749 - loss: 0.3511 - val_accuracy: 0.8740 - val_loss: 0.3526
Epoch 7/100

94/94 1s 9ms/step -
accuracy: 0.8803 - loss: 0.3371 - val_accuracy: 0.8783 - val_loss: 0.3428
Epoch 8/100

94/94 1s 12ms/step -
accuracy: 0.8843 - loss: 0.3248 - val_accuracy: 0.8790 - val_loss: 0.3358
Epoch 9/100

94/94 1s 11ms/step -
accuracy: 0.8874 - loss: 0.3128 - val_accuracy: 0.8813 - val_loss: 0.3331
Epoch 10/100

94/94 1s 12ms/step -
accuracy: 0.8899 - loss: 0.3017 - val_accuracy: 0.8828 - val_loss: 0.3295
Epoch 11/100

94/94 1s 9ms/step -
accuracy: 0.8932 - loss: 0.2932 - val_accuracy: 0.8815 - val_loss: 0.3279
Epoch 12/100

94/94 1s 9ms/step -
accuracy: 0.8960 - loss: 0.2847 - val_accuracy: 0.8828 - val_loss: 0.3252
Epoch 13/100

94/94 1s 11ms/step -
accuracy: 0.8977 - loss: 0.2776 - val_accuracy: 0.8834 - val_loss: 0.3210
Epoch 14/100

94/94 1s 12ms/step -
accuracy: 0.8999 - loss: 0.2693 - val_accuracy: 0.8857 - val_loss: 0.3204
Epoch 15/100

94/94 1s 10ms/step -
accuracy: 0.9032 - loss: 0.2618 - val_accuracy: 0.8873 - val_loss: 0.3185
Epoch 16/100

94/94 1s 11ms/step -
accuracy: 0.9053 - loss: 0.2562 - val_accuracy: 0.8874 - val_loss: 0.3183
Epoch 17/100

```

94/94          1s 11ms/step -
accuracy: 0.9077 - loss: 0.2494 - val_accuracy: 0.8857 - val_loss: 0.3225
Epoch 18/100
94/94          1s 10ms/step -
accuracy: 0.9096 - loss: 0.2456 - val_accuracy: 0.8848 - val_loss: 0.3219
Epoch 19/100
94/94          1s 10ms/step -
accuracy: 0.9136 - loss: 0.2384 - val_accuracy: 0.8855 - val_loss: 0.3265
Epoch 20/100
94/94          1s 10ms/step -
accuracy: 0.9142 - loss: 0.2345 - val_accuracy: 0.8839 - val_loss: 0.3307
Epoch 21/100
94/94          1s 10ms/step -
accuracy: 0.9156 - loss: 0.2296 - val_accuracy: 0.8822 - val_loss: 0.3378
Epoch 22/100
94/94          1s 10ms/step -
accuracy: 0.9170 - loss: 0.2262 - val_accuracy: 0.8808 - val_loss: 0.3466
Epoch 23/100
94/94          1s 10ms/step -
accuracy: 0.9181 - loss: 0.2232 - val_accuracy: 0.8769 - val_loss: 0.3588
Epoch 24/100
94/94          1s 11ms/step -
accuracy: 0.9190 - loss: 0.2223 - val_accuracy: 0.8847 - val_loss: 0.3414
Epoch 25/100
94/94          1s 10ms/step -
accuracy: 0.9210 - loss: 0.2173 - val_accuracy: 0.8848 - val_loss: 0.3428
Epoch 26/100
94/94          1s 10ms/step -
accuracy: 0.9223 - loss: 0.2133 - val_accuracy: 0.8867 - val_loss: 0.3391

```

```

[21]: # Evaluation of the model on the validation set
scores3 = model3.evaluate(X_val, y_val)

print(f"Accuracy for Model 3: {round(scores3[1], 4)}, Loss for Model 3:␣
      ↳{round(scores3[0], 4)}")

print("\n")

```

```

375/375          1s 3ms/step -
accuracy: 0.8874 - loss: 0.3328


```

Accuracy for Model 3: 0.8867, Loss for Model 3: 0.3393

```

[95]: plot_model_history([history2, history3], ['Model 2', 'Model 3'], 'Comparison -␣
      ↳Model 2 vs Model 3')

```



Our third model where we reduced the number of nodes in each layer as well as added an early stopping with a patience level of 5 with 100 epochs does better but it is not comparable as number of epochs are higher

Model 4: Nodes increased to 512 on first layer, added a Dropout value of 50%, added a second hidden layer with 256 nodes, added a Dropout value of 50%, added a third hidden layer with 100 nodes.

```
[22]: model4 = Sequential([
    Input(shape=X_train.shape[1:]),
    Flatten(),
    Rescaling(1./255),
    Dense(512, activation='relu'),
    Dropout(0.5),
    Dense(256, activation='relu'),
    Dropout(0.5),
    Dense(100, activation='relu'),
    Dense(num_classes, activation='softmax')
])
print(model4.summary())

# Compile the model
model4.compile(loss='categorical_crossentropy', optimizer='adam',
               metrics=['accuracy'])
```

Model: "sequential_4"

Layer (type)
Param #

Output Shape

└

flatten_4 (Flatten)	(None, 784)	└
↪ 0		
rescaling_4 (Rescaling)	(None, 784)	└
↪ 0		
dense_13 (Dense)	(None, 512)	└
↪401,920		
dropout (Dropout)	(None, 512)	└
↪ 0		
dense_14 (Dense)	(None, 256)	└
↪131,328		
dropout_1 (Dropout)	(None, 256)	└
↪ 0		
dense_15 (Dense)	(None, 100)	└
↪25,700		
dense_16 (Dense)	(None, 10)	└
↪1,010		

Total params: 559,958 (2.14 MB)

Trainable params: 559,958 (2.14 MB)

Non-trainable params: 0 (0.00 B)

None

```
[23]: history4 = model4.fit(X_train, y_train, validation_data=(X_val, y_val),
    ↪epochs=100, batch_size = 512
    ↪,callbacks=[EarlyStopping(monitor='val_accuracy',
    ↪patience=10)])
```

Epoch 1/100

94/94 5s 27ms/step -

accuracy: 0.5372 - loss: 1.2692 - val_accuracy: 0.8238 - val_loss: 0.4889

Epoch 2/100

94/94 2s 23ms/step -

accuracy: 0.8032 - loss: 0.5466 - val_accuracy: 0.8415 - val_loss: 0.4186

Epoch 3/100
94/94 2s 23ms/step -
accuracy: 0.8294 - loss: 0.4753 - val_accuracy: 0.8597 - val_loss: 0.3760
Epoch 4/100
94/94 2s 23ms/step -
accuracy: 0.8412 - loss: 0.4398 - val_accuracy: 0.8625 - val_loss: 0.3719
Epoch 5/100
94/94 2s 23ms/step -
accuracy: 0.8517 - loss: 0.4163 - val_accuracy: 0.8664 - val_loss: 0.3618
Epoch 6/100
94/94 2s 25ms/step -
accuracy: 0.8553 - loss: 0.4015 - val_accuracy: 0.8705 - val_loss: 0.3426
Epoch 7/100
94/94 2s 25ms/step -
accuracy: 0.8553 - loss: 0.3908 - val_accuracy: 0.8726 - val_loss: 0.3408
Epoch 8/100
94/94 3s 27ms/step -
accuracy: 0.8609 - loss: 0.3747 - val_accuracy: 0.8772 - val_loss: 0.3344
Epoch 9/100
94/94 3s 29ms/step -
accuracy: 0.8672 - loss: 0.3706 - val_accuracy: 0.8754 - val_loss: 0.3409
Epoch 10/100
94/94 3s 28ms/step -
accuracy: 0.8678 - loss: 0.3638 - val_accuracy: 0.8788 - val_loss: 0.3235
Epoch 11/100
94/94 3s 28ms/step -
accuracy: 0.8728 - loss: 0.3501 - val_accuracy: 0.8801 - val_loss: 0.3201
Epoch 12/100
94/94 3s 28ms/step -
accuracy: 0.8722 - loss: 0.3469 - val_accuracy: 0.8837 - val_loss: 0.3150
Epoch 13/100
94/94 2s 25ms/step -
accuracy: 0.8751 - loss: 0.3395 - val_accuracy: 0.8850 - val_loss: 0.3181
Epoch 14/100
94/94 2s 25ms/step -
accuracy: 0.8752 - loss: 0.3417 - val_accuracy: 0.8817 - val_loss: 0.3232
Epoch 15/100
94/94 2s 24ms/step -
accuracy: 0.8795 - loss: 0.3310 - val_accuracy: 0.8825 - val_loss: 0.3205
Epoch 16/100
94/94 2s 25ms/step -
accuracy: 0.8804 - loss: 0.3273 - val_accuracy: 0.8838 - val_loss: 0.3127
Epoch 17/100
94/94 2s 26ms/step -
accuracy: 0.8814 - loss: 0.3227 - val_accuracy: 0.8877 - val_loss: 0.3071
Epoch 18/100
94/94 3s 28ms/step -
accuracy: 0.8810 - loss: 0.3240 - val_accuracy: 0.8868 - val_loss: 0.3074

Epoch 19/100
94/94 2s 25ms/step -
accuracy: 0.8838 - loss: 0.3155 - val_accuracy: 0.8873 - val_loss: 0.3069
Epoch 20/100
94/94 2s 24ms/step -
accuracy: 0.8816 - loss: 0.3158 - val_accuracy: 0.8859 - val_loss: 0.3074
Epoch 21/100
94/94 2s 24ms/step -
accuracy: 0.8835 - loss: 0.3129 - val_accuracy: 0.8868 - val_loss: 0.3094
Epoch 22/100
94/94 2s 23ms/step -
accuracy: 0.8850 - loss: 0.3068 - val_accuracy: 0.8875 - val_loss: 0.3082
Epoch 23/100
94/94 3s 27ms/step -
accuracy: 0.8866 - loss: 0.3017 - val_accuracy: 0.8882 - val_loss: 0.3135
Epoch 24/100
94/94 3s 28ms/step -
accuracy: 0.8879 - loss: 0.3013 - val_accuracy: 0.8907 - val_loss: 0.3020
Epoch 25/100
94/94 3s 28ms/step -
accuracy: 0.8892 - loss: 0.2990 - val_accuracy: 0.8909 - val_loss: 0.2952
Epoch 26/100
94/94 3s 27ms/step -
accuracy: 0.8903 - loss: 0.2931 - val_accuracy: 0.8936 - val_loss: 0.2950
Epoch 27/100
94/94 3s 26ms/step -
accuracy: 0.8905 - loss: 0.2963 - val_accuracy: 0.8896 - val_loss: 0.3036
Epoch 28/100
94/94 2s 24ms/step -
accuracy: 0.8954 - loss: 0.2850 - val_accuracy: 0.8926 - val_loss: 0.2981
Epoch 29/100
94/94 2s 24ms/step -
accuracy: 0.8916 - loss: 0.2881 - val_accuracy: 0.8923 - val_loss: 0.2961
Epoch 30/100
94/94 3s 31ms/step -
accuracy: 0.8955 - loss: 0.2822 - val_accuracy: 0.8928 - val_loss: 0.2974
Epoch 31/100
94/94 3s 36ms/step -
accuracy: 0.8934 - loss: 0.2816 - val_accuracy: 0.8931 - val_loss: 0.2950
Epoch 32/100
94/94 3s 29ms/step -
accuracy: 0.8958 - loss: 0.2794 - val_accuracy: 0.8932 - val_loss: 0.2993
Epoch 33/100
94/94 3s 29ms/step -
accuracy: 0.8965 - loss: 0.2781 - val_accuracy: 0.8936 - val_loss: 0.2936
Epoch 34/100
94/94 3s 30ms/step -
accuracy: 0.8979 - loss: 0.2777 - val_accuracy: 0.8898 - val_loss: 0.3049

```
Epoch 35/100
94/94          2s 23ms/step -
accuracy: 0.8964 - loss: 0.2766 - val_accuracy: 0.8926 - val_loss: 0.2944
Epoch 36/100
94/94          2s 24ms/step -
accuracy: 0.8989 - loss: 0.2761 - val_accuracy: 0.8920 - val_loss: 0.2963
```

```
[24]: # Evaluation of the model on the validation set
scores4 = model4.evaluate(X_val, y_val)


print(f"Accuracy for Model 4: {round(scores4[1], 4)}, Loss for Model 4:
      ↳{round(scores4[0], 4)}")

print("\n")
```

```
375/375          1s 4ms/step -
accuracy: 0.8915 - loss: 0.2966
```

Accuracy for Model 4: 0.892, Loss for Model 4: 0.2971

```
[96]: plot_model_history([history3, history4], ['Model 3', 'Model 4'], 'Comparison -
      ↳Model 3 vs Model 4')
```



Here the 4th model improves on 3rd model slightly when we add a dropout on the first two layers of 50%

Model 5: Nodes increased to 256 on first layer, added a Dropout value of 50%, added a second hidden layer with 100 nodes, added a Dropout value of 50%, added a third hidden layer with 20 nodes.

```
[25]: model5 = Sequential([
    Input(shape=X_train.shape[1:]),
    Flatten(),
    Rescaling(1./255),
    Dense(256, activation='relu'),
    Dropout(0.5),
    Dense(100, activation='relu'),
    Dropout(0.5),
    Dense(20, activation='relu'),
    Dense(num_classes, activation='softmax')
])
print(model5.summary())

# Compile the model
model5.compile(loss='categorical_crossentropy', optimizer='adam',
               metrics=['accuracy'])
```

Model: "sequential_5"

Layer (type) ↳ Param #	Output Shape	
flatten_5 (Flatten) ↳ 0	(None, 784)	↳
rescaling_5 (Rescaling) ↳ 0	(None, 784)	↳
dense_17 (Dense) ↳ 200,960	(None, 256)	↳
dropout_2 (Dropout) ↳ 0	(None, 256)	↳
dense_18 (Dense) ↳ 25,700	(None, 100)	↳
dropout_3 (Dropout) ↳ 0	(None, 100)	↳
dense_19 (Dense) ↳ 2,020	(None, 20)	↳

dense_20 (Dense) (None, 10)
↪210

Total params: 228,890 (894.10 KB)

Trainable params: 228,890 (894.10 KB)

Non-trainable params: 0 (0.00 B)

None

```
[26]: history5 = model5.fit(X_train, y_train, validation_data=(X_val, y_val),  
    ↪epochs=100, batch_size = 512  
    ↪,callbacks=[EarlyStopping(monitor='val_accuracy',  
    ↪patience=10)])
```

Epoch 1/100

94/94 5s 20ms/step -

accuracy: 0.3851 - loss: 1.6822 - val_accuracy: 0.7893 - val_loss: 0.5955

Epoch 2/100

94/94 1s 14ms/step -

accuracy: 0.7534 - loss: 0.7016 - val_accuracy: 0.8225 - val_loss: 0.4785

Epoch 3/100

94/94 2s 16ms/step -

accuracy: 0.7957 - loss: 0.5778 - val_accuracy: 0.8420 - val_loss: 0.4273

Epoch 4/100

94/94 1s 15ms/step -

accuracy: 0.8192 - loss: 0.5173 - val_accuracy: 0.8511 - val_loss: 0.4002

Epoch 5/100

94/94 1s 14ms/step -

accuracy: 0.8298 - loss: 0.4851 - val_accuracy: 0.8553 - val_loss: 0.3911

Epoch 6/100

94/94 2s 16ms/step -

accuracy: 0.8388 - loss: 0.4581 - val_accuracy: 0.8604 - val_loss: 0.3790

Epoch 7/100

94/94 1s 15ms/step -

accuracy: 0.8447 - loss: 0.4409 - val_accuracy: 0.8652 - val_loss: 0.3674

Epoch 8/100

94/94 2s 15ms/step -

accuracy: 0.8496 - loss: 0.4280 - val_accuracy: 0.8694 - val_loss: 0.3577

Epoch 9/100

94/94 2s 19ms/step -

accuracy: 0.8519 - loss: 0.4185 - val_accuracy: 0.8659 - val_loss: 0.3595

Epoch 10/100

94/94 2s 16ms/step -

accuracy: 0.8537 - loss: 0.4082 - val_accuracy: 0.8725 - val_loss: 0.3443
 Epoch 11/100
 94/94 2s 18ms/step -
 accuracy: 0.8585 - loss: 0.4010 - val_accuracy: 0.8714 - val_loss: 0.3453
 Epoch 12/100
 94/94 2s 17ms/step -
 accuracy: 0.8631 - loss: 0.3938 - val_accuracy: 0.8723 - val_loss: 0.3448
 Epoch 13/100
 94/94 2s 19ms/step -
 accuracy: 0.8637 - loss: 0.3872 - val_accuracy: 0.8760 - val_loss: 0.3346
 Epoch 14/100
 94/94 2s 16ms/step -
 accuracy: 0.8690 - loss: 0.3742 - val_accuracy: 0.8783 - val_loss: 0.3329
 Epoch 15/100
 94/94 2s 16ms/step -
 accuracy: 0.8688 - loss: 0.3672 - val_accuracy: 0.8800 - val_loss: 0.3308
 Epoch 16/100
 94/94 1s 15ms/step -
 accuracy: 0.8715 - loss: 0.3609 - val_accuracy: 0.8791 - val_loss: 0.3297
 Epoch 17/100
 94/94 2s 17ms/step -
 accuracy: 0.8717 - loss: 0.3593 - val_accuracy: 0.8786 - val_loss: 0.3250
 Epoch 18/100
 94/94 2s 16ms/step -
 accuracy: 0.8744 - loss: 0.3550 - val_accuracy: 0.8810 - val_loss: 0.3213
 Epoch 19/100
 94/94 2s 16ms/step -
 accuracy: 0.8744 - loss: 0.3527 - val_accuracy: 0.8806 - val_loss: 0.3213
 Epoch 20/100
 94/94 2s 15ms/step -
 accuracy: 0.8755 - loss: 0.3455 - val_accuracy: 0.8805 - val_loss: 0.3200
 Epoch 21/100
 94/94 2s 16ms/step -
 accuracy: 0.8768 - loss: 0.3404 - val_accuracy: 0.8832 - val_loss: 0.3180
 Epoch 22/100
 94/94 1s 15ms/step -
 accuracy: 0.8796 - loss: 0.3375 - val_accuracy: 0.8815 - val_loss: 0.3192
 Epoch 23/100
 94/94 2s 16ms/step -
 accuracy: 0.8784 - loss: 0.3364 - val_accuracy: 0.8845 - val_loss: 0.3194
 Epoch 24/100
 94/94 2s 15ms/step -
 accuracy: 0.8814 - loss: 0.3345 - val_accuracy: 0.8855 - val_loss: 0.3135
 Epoch 25/100
 94/94 2s 15ms/step -
 accuracy: 0.8834 - loss: 0.3272 - val_accuracy: 0.8849 - val_loss: 0.3146
 Epoch 26/100
 94/94 2s 16ms/step -

accuracy: 0.8803 - loss: 0.3292 - val_accuracy: 0.8857 - val_loss: 0.3114
 Epoch 27/100
 94/94 2s 15ms/step -
 accuracy: 0.8825 - loss: 0.3258 - val_accuracy: 0.8866 - val_loss: 0.3073
 Epoch 28/100
 94/94 1s 14ms/step -
 accuracy: 0.8847 - loss: 0.3215 - val_accuracy: 0.8880 - val_loss: 0.3048
 Epoch 29/100
 94/94 1s 14ms/step -
 accuracy: 0.8832 - loss: 0.3162 - val_accuracy: 0.8860 - val_loss: 0.3109
 Epoch 30/100
 94/94 2s 15ms/step -
 accuracy: 0.8869 - loss: 0.3131 - val_accuracy: 0.8889 - val_loss: 0.3094
 Epoch 31/100
 94/94 1s 15ms/step -
 accuracy: 0.8865 - loss: 0.3133 - val_accuracy: 0.8873 - val_loss: 0.3036
 Epoch 32/100
 94/94 2s 19ms/step -
 accuracy: 0.8847 - loss: 0.3135 - val_accuracy: 0.8860 - val_loss: 0.3075
 Epoch 33/100
 94/94 1s 14ms/step -
 accuracy: 0.8870 - loss: 0.3092 - val_accuracy: 0.8877 - val_loss: 0.3083
 Epoch 34/100
 94/94 2s 17ms/step -
 accuracy: 0.8888 - loss: 0.3080 - val_accuracy: 0.8923 - val_loss: 0.3032
 Epoch 35/100
 94/94 2s 17ms/step -
 accuracy: 0.8912 - loss: 0.3060 - val_accuracy: 0.8882 - val_loss: 0.3039
 Epoch 36/100
 94/94 2s 16ms/step -
 accuracy: 0.8888 - loss: 0.3072 - val_accuracy: 0.8886 - val_loss: 0.3071
 Epoch 37/100
 94/94 1s 15ms/step -
 accuracy: 0.8907 - loss: 0.3023 - val_accuracy: 0.8897 - val_loss: 0.3066
 Epoch 38/100
 94/94 2s 22ms/step -
 accuracy: 0.8913 - loss: 0.3016 - val_accuracy: 0.8906 - val_loss: 0.3038
 Epoch 39/100
 94/94 2s 21ms/step -
 accuracy: 0.8897 - loss: 0.3021 - val_accuracy: 0.8923 - val_loss: 0.3000
 Epoch 40/100
 94/94 2s 16ms/step -
 accuracy: 0.8916 - loss: 0.2985 - val_accuracy: 0.8920 - val_loss: 0.2981
 Epoch 41/100
 94/94 2s 16ms/step -
 accuracy: 0.8936 - loss: 0.2942 - val_accuracy: 0.8929 - val_loss: 0.3004
 Epoch 42/100
 94/94 2s 15ms/step -

```

accuracy: 0.8908 - loss: 0.2942 - val_accuracy: 0.8945 - val_loss: 0.2986
Epoch 43/100
94/94          2s 16ms/step -
accuracy: 0.8932 - loss: 0.2932 - val_accuracy: 0.8924 - val_loss: 0.2965
Epoch 44/100
94/94          2s 16ms/step -
accuracy: 0.8930 - loss: 0.2944 - val_accuracy: 0.8940 - val_loss: 0.2962
Epoch 45/100
94/94          1s 15ms/step -
accuracy: 0.8964 - loss: 0.2825 - val_accuracy: 0.8916 - val_loss: 0.2992
Epoch 46/100
94/94          2s 16ms/step -
accuracy: 0.8957 - loss: 0.2877 - val_accuracy: 0.8904 - val_loss: 0.3037
Epoch 47/100
94/94          1s 15ms/step -
accuracy: 0.8930 - loss: 0.2901 - val_accuracy: 0.8917 - val_loss: 0.2970
Epoch 48/100
94/94          2s 17ms/step -
accuracy: 0.8966 - loss: 0.2834 - val_accuracy: 0.8904 - val_loss: 0.3068
Epoch 49/100
94/94          1s 15ms/step -
accuracy: 0.8946 - loss: 0.2869 - val_accuracy: 0.8930 - val_loss: 0.3005
Epoch 50/100
94/94          1s 15ms/step -
accuracy: 0.8998 - loss: 0.2769 - val_accuracy: 0.8888 - val_loss: 0.3050
Epoch 51/100
94/94          1s 15ms/step -
accuracy: 0.8972 - loss: 0.2820 - val_accuracy: 0.8894 - val_loss: 0.3074
Epoch 52/100
94/94          2s 16ms/step -
accuracy: 0.8983 - loss: 0.2813 - val_accuracy: 0.8916 - val_loss: 0.3034

```

[27]: *# Evaluation of the model on the validation set*

```

scores5 = model5.evaluate(X_val, y_val)

print("\n")
print(f"Accuracy for Model 5: {round(scores5[1], 4)}, Loss for Model 5:  
     {round(scores5[0], 4)}")

```


```

375/375          1s 3ms/step -
accuracy: 0.8881 - loss: 0.3049

```

Accuracy for Model 5: 0.8916, Loss for Model 5: 0.3041

[98]: `plot_model_history([history4, history5], ['Model 4', 'Model 5'], 'Comparison -  
  Model 4 vs Model 5')`



The 5th model is the same as 4th, it only differs in the number of nodes, where we reduce them. As you can see it does not improve the accuracy or the loss

0.8 5. Try to improve the accuracy of your model by using convolution.

0.8.1 Train at least two different models (you can vary the number of convolutional and pooling layers or whether you include a fully connected layer before the output, etc.)

```
[28]: from keras.layers import Reshape

preprocess = Sequential([
    Reshape(target_shape=(X_train.shape[1], X_train.shape[2], 1)), #_
    ↪explicitly state the 4th (channel) dimension
    Rescaling(1./255)
])

X_sets = [X_train, X_test, X_val]
X_train_4D, X_test_4D, X_val_4D = [preprocess(X) for X in X_sets]
```

Model 6: Adding Convolution, Pooling, and Fully Connected Layer with No Hidden Layer

```
[29]: from keras.layers import Conv2D, MaxPooling2D

# Build the model
model6 = Sequential([
    Input(shape=X_train_4D.shape[1:]),
    Conv2D(32, (3, 3), activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),
```

```

        Flatten(),
        Dense(num_classes, activation='softmax')
    ])

    # Compile the model
    model6.compile(loss='categorical_crossentropy', optimizer='adam',
        metrics=['accuracy'])
    print(model6.summary())

```

Model: "sequential_7"

Layer (type)	Output Shape	
Param #		
conv2d (Conv2D)	(None, 26, 26, 32)	
320		
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	
0		
flatten_6 (Flatten)	(None, 5408)	
0		
dense_21 (Dense)	(None, 10)	
54,090		

Total params: 54,410 (212.54 KB)

Trainable params: 54,410 (212.54 KB)

Non-trainable params: 0 (0.00 B)

None

```

[30]: # Fit the model
history6 = model6.fit(
    X_train_4D, y_train, validation_data=(X_val_4D, y_val), epochs=100,
    batch_size=2048,
    callbacks=[EarlyStopping(monitor='val_accuracy', patience=10)]
)

```

Epoch 1/100

24/24 5s 160ms/step -
 accuracy: 0.4792 - loss: 1.7679 - val_accuracy: 0.7388 - val_loss: 0.8033
 Epoch 2/100
 24/24 3s 143ms/step -
 accuracy: 0.7583 - loss: 0.7315 - val_accuracy: 0.7792 - val_loss: 0.6113
 Epoch 3/100
 24/24 4s 157ms/step -
 accuracy: 0.7979 - loss: 0.5844 - val_accuracy: 0.8102 - val_loss: 0.5344
 Epoch 4/100
 24/24 4s 150ms/step -
 accuracy: 0.8251 - loss: 0.5165 - val_accuracy: 0.8267 - val_loss: 0.4903
 Epoch 5/100
 24/24 4s 156ms/step -
 accuracy: 0.8380 - loss: 0.4740 - val_accuracy: 0.8378 - val_loss: 0.4633
 Epoch 6/100
 24/24 4s 151ms/step -
 accuracy: 0.8484 - loss: 0.4457 - val_accuracy: 0.8444 - val_loss: 0.4432
 Epoch 7/100
 24/24 4s 146ms/step -
 accuracy: 0.8561 - loss: 0.4250 - val_accuracy: 0.8501 - val_loss: 0.4262
 Epoch 8/100
 24/24 4s 146ms/step -
 accuracy: 0.8620 - loss: 0.4083 - val_accuracy: 0.8558 - val_loss: 0.4118
 Epoch 9/100
 24/24 4s 146ms/step -
 accuracy: 0.8672 - loss: 0.3944 - val_accuracy: 0.8601 - val_loss: 0.4001
 Epoch 10/100
 24/24 4s 150ms/step -
 accuracy: 0.8710 - loss: 0.3827 - val_accuracy: 0.8636 - val_loss: 0.3905
 Epoch 11/100
 24/24 4s 160ms/step -
 accuracy: 0.8746 - loss: 0.3727 - val_accuracy: 0.8657 - val_loss: 0.3822
 Epoch 12/100
 24/24 4s 147ms/step -
 accuracy: 0.8771 - loss: 0.3641 - val_accuracy: 0.8689 - val_loss: 0.3749
 Epoch 13/100
 24/24 4s 147ms/step -
 accuracy: 0.8801 - loss: 0.3564 - val_accuracy: 0.8707 - val_loss: 0.3684
 Epoch 14/100
 24/24 4s 146ms/step -
 accuracy: 0.8813 - loss: 0.3494 - val_accuracy: 0.8724 - val_loss: 0.3625
 Epoch 15/100
 24/24 4s 149ms/step -
 accuracy: 0.8830 - loss: 0.3431 - val_accuracy: 0.8742 - val_loss: 0.3572
 Epoch 16/100
 24/24 4s 154ms/step -
 accuracy: 0.8843 - loss: 0.3373 - val_accuracy: 0.8756 - val_loss: 0.3523
 Epoch 17/100

24/24 4s 161ms/step -
 accuracy: 0.8865 - loss: 0.3319 - val_accuracy: 0.8777 - val_loss: 0.3478
 Epoch 18/100
 24/24 4s 158ms/step -
 accuracy: 0.8881 - loss: 0.3269 - val_accuracy: 0.8788 - val_loss: 0.3436
 Epoch 19/100
 24/24 4s 146ms/step -
 accuracy: 0.8898 - loss: 0.3222 - val_accuracy: 0.8801 - val_loss: 0.3398
 Epoch 20/100
 24/24 4s 147ms/step -
 accuracy: 0.8914 - loss: 0.3179 - val_accuracy: 0.8815 - val_loss: 0.3363
 Epoch 21/100
 24/24 4s 144ms/step -
 accuracy: 0.8924 - loss: 0.3138 - val_accuracy: 0.8821 - val_loss: 0.3330
 Epoch 22/100
 24/24 4s 154ms/step -
 accuracy: 0.8932 - loss: 0.3099 - val_accuracy: 0.8829 - val_loss: 0.3299
 Epoch 23/100
 24/24 4s 160ms/step -
 accuracy: 0.8945 - loss: 0.3063 - val_accuracy: 0.8836 - val_loss: 0.3270
 Epoch 24/100
 24/24 4s 168ms/step -
 accuracy: 0.8956 - loss: 0.3028 - val_accuracy: 0.8849 - val_loss: 0.3244
 Epoch 25/100
 24/24 4s 150ms/step -
 accuracy: 0.8965 - loss: 0.2996 - val_accuracy: 0.8857 - val_loss: 0.3218
 Epoch 26/100
 24/24 4s 159ms/step -
 accuracy: 0.8976 - loss: 0.2965 - val_accuracy: 0.8866 - val_loss: 0.3195
 Epoch 27/100
 24/24 4s 148ms/step -
 accuracy: 0.8985 - loss: 0.2935 - val_accuracy: 0.8874 - val_loss: 0.3173
 Epoch 28/100
 24/24 4s 158ms/step -
 accuracy: 0.8995 - loss: 0.2907 - val_accuracy: 0.8887 - val_loss: 0.3152
 Epoch 29/100
 24/24 4s 175ms/step -
 accuracy: 0.9003 - loss: 0.2880 - val_accuracy: 0.8892 - val_loss: 0.3132
 Epoch 30/100
 24/24 4s 162ms/step -
 accuracy: 0.9010 - loss: 0.2854 - val_accuracy: 0.8904 - val_loss: 0.3114
 Epoch 31/100
 24/24 4s 184ms/step -
 accuracy: 0.9016 - loss: 0.2829 - val_accuracy: 0.8908 - val_loss: 0.3096
 Epoch 32/100
 24/24 4s 154ms/step -
 accuracy: 0.9020 - loss: 0.2806 - val_accuracy: 0.8911 - val_loss: 0.3080
 Epoch 33/100

24/24 4s 167ms/step -
 accuracy: 0.9028 - loss: 0.2783 - val_accuracy: 0.8916 - val_loss: 0.3064
 Epoch 34/100
 24/24 4s 157ms/step -
 accuracy: 0.9037 - loss: 0.2761 - val_accuracy: 0.8917 - val_loss: 0.3049
 Epoch 35/100
 24/24 4s 157ms/step -
 accuracy: 0.9049 - loss: 0.2740 - val_accuracy: 0.8922 - val_loss: 0.3035
 Epoch 36/100
 24/24 4s 153ms/step -
 accuracy: 0.9055 - loss: 0.2720 - val_accuracy: 0.8927 - val_loss: 0.3022
 Epoch 37/100
 24/24 4s 148ms/step -
 accuracy: 0.9063 - loss: 0.2700 - val_accuracy: 0.8932 - val_loss: 0.3009
 Epoch 38/100
 24/24 4s 148ms/step -
 accuracy: 0.9072 - loss: 0.2681 - val_accuracy: 0.8929 - val_loss: 0.2997
 Epoch 39/100
 24/24 4s 150ms/step -
 accuracy: 0.9077 - loss: 0.2662 - val_accuracy: 0.8932 - val_loss: 0.2985
 Epoch 40/100
 24/24 4s 156ms/step -
 accuracy: 0.9084 - loss: 0.2644 - val_accuracy: 0.8936 - val_loss: 0.2974
 Epoch 41/100
 24/24 4s 152ms/step -
 accuracy: 0.9089 - loss: 0.2627 - val_accuracy: 0.8939 - val_loss: 0.2963
 Epoch 42/100
 24/24 4s 158ms/step -
 accuracy: 0.9095 - loss: 0.2610 - val_accuracy: 0.8938 - val_loss: 0.2952
 Epoch 43/100
 24/24 4s 154ms/step -
 accuracy: 0.9102 - loss: 0.2593 - val_accuracy: 0.8940 - val_loss: 0.2942
 Epoch 44/100
 24/24 4s 160ms/step -
 accuracy: 0.9108 - loss: 0.2577 - val_accuracy: 0.8946 - val_loss: 0.2933
 Epoch 45/100
 24/24 4s 148ms/step -
 accuracy: 0.9113 - loss: 0.2562 - val_accuracy: 0.8948 - val_loss: 0.2924
 Epoch 46/100
 24/24 4s 150ms/step -
 accuracy: 0.9118 - loss: 0.2547 - val_accuracy: 0.8951 - val_loss: 0.2915
 Epoch 47/100
 24/24 4s 150ms/step -
 accuracy: 0.9126 - loss: 0.2532 - val_accuracy: 0.8957 - val_loss: 0.2906
 Epoch 48/100
 24/24 4s 159ms/step -
 accuracy: 0.9128 - loss: 0.2517 - val_accuracy: 0.8959 - val_loss: 0.2898
 Epoch 49/100

24/24 4s 153ms/step -
 accuracy: 0.9135 - loss: 0.2503 - val_accuracy: 0.8960 - val_loss: 0.2890
 Epoch 50/100
 24/24 4s 154ms/step -
 accuracy: 0.9138 - loss: 0.2489 - val_accuracy: 0.8962 - val_loss: 0.2883
 Epoch 51/100
 24/24 4s 156ms/step -
 accuracy: 0.9142 - loss: 0.2476 - val_accuracy: 0.8959 - val_loss: 0.2876
 Epoch 52/100
 24/24 4s 149ms/step -
 accuracy: 0.9147 - loss: 0.2462 - val_accuracy: 0.8958 - val_loss: 0.2869
 Epoch 53/100
 24/24 4s 149ms/step -
 accuracy: 0.9152 - loss: 0.2450 - val_accuracy: 0.8961 - val_loss: 0.2862
 Epoch 54/100
 24/24 4s 155ms/step -
 accuracy: 0.9158 - loss: 0.2437 - val_accuracy: 0.8966 - val_loss: 0.2856
 Epoch 55/100
 24/24 4s 145ms/step -
 accuracy: 0.9162 - loss: 0.2424 - val_accuracy: 0.8967 - val_loss: 0.2850
 Epoch 56/100
 24/24 4s 147ms/step -
 accuracy: 0.9167 - loss: 0.2412 - val_accuracy: 0.8967 - val_loss: 0.2844
 Epoch 57/100
 24/24 4s 158ms/step -
 accuracy: 0.9174 - loss: 0.2400 - val_accuracy: 0.8966 - val_loss: 0.2838
 Epoch 58/100
 24/24 4s 149ms/step -
 accuracy: 0.9177 - loss: 0.2388 - val_accuracy: 0.8968 - val_loss: 0.2833
 Epoch 59/100
 24/24 4s 149ms/step -
 accuracy: 0.9181 - loss: 0.2377 - val_accuracy: 0.8972 - val_loss: 0.2828
 Epoch 60/100
 24/24 4s 148ms/step -
 accuracy: 0.9187 - loss: 0.2366 - val_accuracy: 0.8973 - val_loss: 0.2823
 Epoch 61/100
 24/24 4s 158ms/step -
 accuracy: 0.9190 - loss: 0.2354 - val_accuracy: 0.8973 - val_loss: 0.2818
 Epoch 62/100
 24/24 4s 150ms/step -
 accuracy: 0.9192 - loss: 0.2343 - val_accuracy: 0.8978 - val_loss: 0.2812
 Epoch 63/100
 24/24 4s 152ms/step -
 accuracy: 0.9196 - loss: 0.2332 - val_accuracy: 0.8979 - val_loss: 0.2807
 Epoch 64/100
 24/24 4s 149ms/step -
 accuracy: 0.9200 - loss: 0.2321 - val_accuracy: 0.8980 - val_loss: 0.2803
 Epoch 65/100

24/24 4s 150ms/step -
 accuracy: 0.9205 - loss: 0.2311 - val_accuracy: 0.8982 - val_loss: 0.2798
 Epoch 66/100
 24/24 4s 152ms/step -
 accuracy: 0.9210 - loss: 0.2300 - val_accuracy: 0.8982 - val_loss: 0.2793
 Epoch 67/100
 24/24 4s 149ms/step -
 accuracy: 0.9213 - loss: 0.2290 - val_accuracy: 0.8985 - val_loss: 0.2788
 Epoch 68/100
 24/24 4s 147ms/step -
 accuracy: 0.9220 - loss: 0.2279 - val_accuracy: 0.8990 - val_loss: 0.2784
 Epoch 69/100
 24/24 4s 146ms/step -
 accuracy: 0.9225 - loss: 0.2269 - val_accuracy: 0.8989 - val_loss: 0.2779
 Epoch 70/100
 24/24 4s 187ms/step -
 accuracy: 0.9227 - loss: 0.2259 - val_accuracy: 0.8989 - val_loss: 0.2774
 Epoch 71/100
 24/24 4s 153ms/step -
 accuracy: 0.9232 - loss: 0.2249 - val_accuracy: 0.8989 - val_loss: 0.2770
 Epoch 72/100
 24/24 4s 153ms/step -
 accuracy: 0.9235 - loss: 0.2239 - val_accuracy: 0.8992 - val_loss: 0.2766
 Epoch 73/100
 24/24 4s 156ms/step -
 accuracy: 0.9241 - loss: 0.2230 - val_accuracy: 0.8993 - val_loss: 0.2761
 Epoch 74/100
 24/24 4s 153ms/step -
 accuracy: 0.9243 - loss: 0.2220 - val_accuracy: 0.8996 - val_loss: 0.2757
 Epoch 75/100
 24/24 4s 154ms/step -
 accuracy: 0.9245 - loss: 0.2210 - val_accuracy: 0.8995 - val_loss: 0.2753
 Epoch 76/100
 24/24 4s 155ms/step -
 accuracy: 0.9250 - loss: 0.2201 - val_accuracy: 0.8998 - val_loss: 0.2749
 Epoch 77/100
 24/24 4s 151ms/step -
 accuracy: 0.9253 - loss: 0.2192 - val_accuracy: 0.8997 - val_loss: 0.2746
 Epoch 78/100
 24/24 4s 152ms/step -
 accuracy: 0.9257 - loss: 0.2183 - val_accuracy: 0.9003 - val_loss: 0.2742
 Epoch 79/100
 24/24 4s 153ms/step -
 accuracy: 0.9262 - loss: 0.2174 - val_accuracy: 0.9001 - val_loss: 0.2738
 Epoch 80/100
 24/24 4s 153ms/step -
 accuracy: 0.9263 - loss: 0.2165 - val_accuracy: 0.9004 - val_loss: 0.2735
 Epoch 81/100

24/24 4s 152ms/step -
 accuracy: 0.9267 - loss: 0.2156 - val_accuracy: 0.9007 - val_loss: 0.2732
 Epoch 82/100
 24/24 4s 153ms/step -
 accuracy: 0.9269 - loss: 0.2148 - val_accuracy: 0.9012 - val_loss: 0.2729
 Epoch 83/100
 24/24 4s 154ms/step -
 accuracy: 0.9271 - loss: 0.2139 - val_accuracy: 0.9013 - val_loss: 0.2726
 Epoch 84/100
 24/24 4s 152ms/step -
 accuracy: 0.9272 - loss: 0.2131 - val_accuracy: 0.9014 - val_loss: 0.2723
 Epoch 85/100
 24/24 4s 157ms/step -
 accuracy: 0.9275 - loss: 0.2123 - val_accuracy: 0.9018 - val_loss: 0.2721
 Epoch 86/100
 24/24 4s 152ms/step -
 accuracy: 0.9278 - loss: 0.2115 - val_accuracy: 0.9020 - val_loss: 0.2718
 Epoch 87/100
 24/24 4s 156ms/step -
 accuracy: 0.9282 - loss: 0.2107 - val_accuracy: 0.9020 - val_loss: 0.2716
 Epoch 88/100
 24/24 4s 154ms/step -
 accuracy: 0.9283 - loss: 0.2099 - val_accuracy: 0.9023 - val_loss: 0.2713
 Epoch 89/100
 24/24 4s 152ms/step -
 accuracy: 0.9289 - loss: 0.2091 - val_accuracy: 0.9023 - val_loss: 0.2711
 Epoch 90/100
 24/24 4s 163ms/step -
 accuracy: 0.9293 - loss: 0.2083 - val_accuracy: 0.9024 - val_loss: 0.2709
 Epoch 91/100
 24/24 5s 203ms/step -
 accuracy: 0.9295 - loss: 0.2076 - val_accuracy: 0.9025 - val_loss: 0.2707
 Epoch 92/100
 24/24 4s 173ms/step -
 accuracy: 0.9297 - loss: 0.2068 - val_accuracy: 0.9027 - val_loss: 0.2705
 Epoch 93/100
 24/24 4s 159ms/step -
 accuracy: 0.9303 - loss: 0.2060 - val_accuracy: 0.9029 - val_loss: 0.2703
 Epoch 94/100
 24/24 4s 159ms/step -
 accuracy: 0.9308 - loss: 0.2053 - val_accuracy: 0.9030 - val_loss: 0.2701
 Epoch 95/100
 24/24 4s 152ms/step -
 accuracy: 0.9309 - loss: 0.2046 - val_accuracy: 0.9034 - val_loss: 0.2699
 Epoch 96/100
 24/24 4s 155ms/step -
 accuracy: 0.9312 - loss: 0.2038 - val_accuracy: 0.9038 - val_loss: 0.2698
 Epoch 97/100

```

24/24          4s 156ms/step -
accuracy: 0.9312 - loss: 0.2031 - val_accuracy: 0.9039 - val_loss: 0.2696
Epoch 98/100
24/24          4s 174ms/step -
accuracy: 0.9314 - loss: 0.2024 - val_accuracy: 0.9039 - val_loss: 0.2694
Epoch 99/100
24/24          4s 172ms/step -
accuracy: 0.9317 - loss: 0.2017 - val_accuracy: 0.9042 - val_loss: 0.2693
Epoch 100/100
24/24          4s 176ms/step -
accuracy: 0.9321 - loss: 0.2010 - val_accuracy: 0.9043 - val_loss: 0.2691

```

```

[31]: # Evaluation of the model on the validation set
scores6 = model6.evaluate(X_val_4D, y_val)

print(f"Accuracy for Model: {round(scores6[1], 4)}, Loss for Model: ␣
      ↪{round(scores6[0], 4)}")

print("\n")

```

```

375/375        2s 5ms/step -
accuracy: 0.9027 - loss: 0.2692
Accuracy for Model: 0.9043, Loss for Model: 0.2692

```

Model 7: Adding Convolution, Pooling, and Fully Connected Layer with a hidden layer (100 Nodes)

```

[32]: model7 = Sequential([
    Input(shape=X_train_4D.shape[1:]),
    Conv2D(32, (3, 3), activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),
    Flatten(),
    Dense(100, activation='relu'),
    Dense(num_classes, activation='softmax')
])

# Compile the model
model7.compile(loss='categorical_crossentropy', optimizer='adam', ␣
              ↪metrics=['accuracy'])
print(model7.summary())

```

Model: "sequential_8"

Layer (type)
↪Param #

Output Shape

␣

conv2d_1 (Conv2D)	(None, 26, 26, 32)	└
↪ 320		
max_pooling2d_1 (MaxPooling2D)	(None, 13, 13, 32)	└
↪ 0		
flatten_7 (Flatten)	(None, 5408)	└
↪ 0		
dense_22 (Dense)	(None, 100)	└
↪ 540,900		
dense_23 (Dense)	(None, 10)	└
↪ 1,010		

Total params: 542,230 (2.07 MB)

Trainable params: 542,230 (2.07 MB)

Non-trainable params: 0 (0.00 B)

None

```
[33]: # Fit the model
history7 = model7.fit(
    X_train_4D, y_train, validation_data=(X_val_4D, y_val), epochs=100,
    ↪ batch_size=2048,
    callbacks=[EarlyStopping(monitor='val_accuracy', patience=5)]
)
```

Epoch 1/100

24/24 7s 202ms/step -

accuracy: 0.5619 - loss: 1.4860 - val_accuracy: 0.7822 - val_loss: 0.6049

Epoch 2/100

24/24 5s 188ms/step -

accuracy: 0.7997 - loss: 0.5637 - val_accuracy: 0.8307 - val_loss: 0.4818

Epoch 3/100

24/24 5s 196ms/step -

accuracy: 0.8411 - loss: 0.4618 - val_accuracy: 0.8477 - val_loss: 0.4290

Epoch 4/100

24/24 5s 200ms/step -

accuracy: 0.8554 - loss: 0.4183 - val_accuracy: 0.8562 - val_loss: 0.4070

Epoch 5/100

24/24 5s 193ms/step -
 accuracy: 0.8636 - loss: 0.3931 - val_accuracy: 0.8654 - val_loss: 0.3856
 Epoch 6/100
 24/24 5s 193ms/step -
 accuracy: 0.8731 - loss: 0.3684 - val_accuracy: 0.8709 - val_loss: 0.3710
 Epoch 7/100
 24/24 5s 187ms/step -
 accuracy: 0.8792 - loss: 0.3495 - val_accuracy: 0.8754 - val_loss: 0.3564
 Epoch 8/100
 24/24 5s 187ms/step -
 accuracy: 0.8840 - loss: 0.3349 - val_accuracy: 0.8784 - val_loss: 0.3468
 Epoch 9/100
 24/24 5s 186ms/step -
 accuracy: 0.8881 - loss: 0.3242 - val_accuracy: 0.8810 - val_loss: 0.3391
 Epoch 10/100
 24/24 5s 186ms/step -
 accuracy: 0.8911 - loss: 0.3148 - val_accuracy: 0.8838 - val_loss: 0.3319
 Epoch 11/100
 24/24 5s 194ms/step -
 accuracy: 0.8938 - loss: 0.3054 - val_accuracy: 0.8857 - val_loss: 0.3247
 Epoch 12/100
 24/24 5s 201ms/step -
 accuracy: 0.8968 - loss: 0.2966 - val_accuracy: 0.8875 - val_loss: 0.3183
 Epoch 13/100
 24/24 4s 184ms/step -
 accuracy: 0.8995 - loss: 0.2889 - val_accuracy: 0.8893 - val_loss: 0.3131
 Epoch 14/100
 24/24 5s 207ms/step -
 accuracy: 0.9014 - loss: 0.2823 - val_accuracy: 0.8909 - val_loss: 0.3086
 Epoch 15/100
 24/24 5s 187ms/step -
 accuracy: 0.9032 - loss: 0.2762 - val_accuracy: 0.8920 - val_loss: 0.3045
 Epoch 16/100
 24/24 5s 199ms/step -
 accuracy: 0.9053 - loss: 0.2703 - val_accuracy: 0.8939 - val_loss: 0.3005
 Epoch 17/100
 24/24 5s 209ms/step -
 accuracy: 0.9070 - loss: 0.2646 - val_accuracy: 0.8942 - val_loss: 0.2968
 Epoch 18/100
 24/24 5s 222ms/step -
 accuracy: 0.9088 - loss: 0.2594 - val_accuracy: 0.8971 - val_loss: 0.2916
 Epoch 19/100
 24/24 5s 218ms/step -
 accuracy: 0.9105 - loss: 0.2540 - val_accuracy: 0.8964 - val_loss: 0.2907
 Epoch 20/100
 24/24 5s 214ms/step -
 accuracy: 0.9120 - loss: 0.2496 - val_accuracy: 0.8972 - val_loss: 0.2889
 Epoch 21/100

24/24 5s 193ms/step -
 accuracy: 0.9135 - loss: 0.2454 - val_accuracy: 0.8968 - val_loss: 0.2867
 Epoch 22/100
 24/24 5s 201ms/step -
 accuracy: 0.9148 - loss: 0.2411 - val_accuracy: 0.8971 - val_loss: 0.2850
 Epoch 23/100
 24/24 5s 188ms/step -
 accuracy: 0.9163 - loss: 0.2372 - val_accuracy: 0.8985 - val_loss: 0.2830
 Epoch 24/100
 24/24 5s 188ms/step -
 accuracy: 0.9181 - loss: 0.2331 - val_accuracy: 0.8994 - val_loss: 0.2813
 Epoch 25/100
 24/24 5s 194ms/step -
 accuracy: 0.9201 - loss: 0.2293 - val_accuracy: 0.8999 - val_loss: 0.2800
 Epoch 26/100
 24/24 5s 193ms/step -
 accuracy: 0.9216 - loss: 0.2258 - val_accuracy: 0.9014 - val_loss: 0.2792
 Epoch 27/100
 24/24 5s 189ms/step -
 accuracy: 0.9232 - loss: 0.2226 - val_accuracy: 0.9011 - val_loss: 0.2784
 Epoch 28/100
 24/24 5s 204ms/step -
 accuracy: 0.9239 - loss: 0.2196 - val_accuracy: 0.9022 - val_loss: 0.2784
 Epoch 29/100
 24/24 5s 212ms/step -
 accuracy: 0.9248 - loss: 0.2170 - val_accuracy: 0.9013 - val_loss: 0.2786
 Epoch 30/100
 24/24 5s 191ms/step -
 accuracy: 0.9261 - loss: 0.2147 - val_accuracy: 0.9018 - val_loss: 0.2785
 Epoch 31/100
 24/24 5s 190ms/step -
 accuracy: 0.9271 - loss: 0.2127 - val_accuracy: 0.9017 - val_loss: 0.2753
 Epoch 32/100
 24/24 5s 190ms/step -
 accuracy: 0.9274 - loss: 0.2110 - val_accuracy: 0.9038 - val_loss: 0.2683
 Epoch 33/100
 24/24 5s 189ms/step -
 accuracy: 0.9285 - loss: 0.2076 - val_accuracy: 0.9046 - val_loss: 0.2640
 Epoch 34/100
 24/24 5s 208ms/step -
 accuracy: 0.9307 - loss: 0.2030 - val_accuracy: 0.9048 - val_loss: 0.2636
 Epoch 35/100
 24/24 5s 201ms/step -
 accuracy: 0.9326 - loss: 0.1983 - val_accuracy: 0.9046 - val_loss: 0.2630
 Epoch 36/100
 24/24 5s 200ms/step -
 accuracy: 0.9334 - loss: 0.1950 - val_accuracy: 0.9049 - val_loss: 0.2625
 Epoch 37/100

24/24 5s 196ms/step -
 accuracy: 0.9349 - loss: 0.1923 - val_accuracy: 0.9057 - val_loss: 0.2620
 Epoch 38/100
 24/24 5s 216ms/step -
 accuracy: 0.9362 - loss: 0.1899 - val_accuracy: 0.9059 - val_loss: 0.2617
 Epoch 39/100
 24/24 5s 194ms/step -
 accuracy: 0.9367 - loss: 0.1875 - val_accuracy: 0.9061 - val_loss: 0.2612
 Epoch 40/100
 24/24 5s 189ms/step -
 accuracy: 0.9378 - loss: 0.1849 - val_accuracy: 0.9062 - val_loss: 0.2608
 Epoch 41/100
 24/24 5s 190ms/step -
 accuracy: 0.9392 - loss: 0.1823 - val_accuracy: 0.9066 - val_loss: 0.2603
 Epoch 42/100
 24/24 5s 197ms/step -
 accuracy: 0.9401 - loss: 0.1798 - val_accuracy: 0.9068 - val_loss: 0.2600
 Epoch 43/100
 24/24 5s 223ms/step -
 accuracy: 0.9407 - loss: 0.1773 - val_accuracy: 0.9073 - val_loss: 0.2597
 Epoch 44/100
 24/24 5s 225ms/step -
 accuracy: 0.9415 - loss: 0.1749 - val_accuracy: 0.9076 - val_loss: 0.2594
 Epoch 45/100
 24/24 6s 239ms/step -
 accuracy: 0.9426 - loss: 0.1726 - val_accuracy: 0.9081 - val_loss: 0.2589
 Epoch 46/100
 24/24 6s 229ms/step -
 accuracy: 0.9432 - loss: 0.1702 - val_accuracy: 0.9082 - val_loss: 0.2588
 Epoch 47/100
 24/24 5s 186ms/step -
 accuracy: 0.9440 - loss: 0.1680 - val_accuracy: 0.9087 - val_loss: 0.2588
 Epoch 48/100
 24/24 5s 209ms/step -
 accuracy: 0.9444 - loss: 0.1658 - val_accuracy: 0.9088 - val_loss: 0.2586
 Epoch 49/100
 24/24 5s 219ms/step -
 accuracy: 0.9454 - loss: 0.1635 - val_accuracy: 0.9087 - val_loss: 0.2584
 Epoch 50/100
 24/24 5s 216ms/step -
 accuracy: 0.9461 - loss: 0.1614 - val_accuracy: 0.9095 - val_loss: 0.2586
 Epoch 51/100
 24/24 5s 193ms/step -
 accuracy: 0.9469 - loss: 0.1594 - val_accuracy: 0.9106 - val_loss: 0.2587
 Epoch 52/100
 24/24 5s 191ms/step -
 accuracy: 0.9478 - loss: 0.1573 - val_accuracy: 0.9107 - val_loss: 0.2587
 Epoch 53/100

```

24/24          5s 199ms/step -
accuracy: 0.9484 - loss: 0.1553 - val_accuracy: 0.9101 - val_loss: 0.2589
Epoch 54/100
24/24          5s 209ms/step -
accuracy: 0.9492 - loss: 0.1533 - val_accuracy: 0.9103 - val_loss: 0.2590
Epoch 55/100
24/24          5s 200ms/step -
accuracy: 0.9498 - loss: 0.1513 - val_accuracy: 0.9102 - val_loss: 0.2587
Epoch 56/100
24/24          5s 201ms/step -
accuracy: 0.9507 - loss: 0.1491 - val_accuracy: 0.9103 - val_loss: 0.2590
Epoch 57/100
24/24          5s 193ms/step -
accuracy: 0.9517 - loss: 0.1470 - val_accuracy: 0.9101 - val_loss: 0.2588

```

```

[34]: # Evaluation of the model on the validation set
scores7 = model7.evaluate(X_val_4D, y_val)
print(f"Accuracy for Model: {round(scores7[1], 4)}, Loss for Model: {round(scores7[0], 4)}")
print("\n")

```

```

375/375        2s 5ms/step -
accuracy: 0.9120 - loss: 0.2546
Accuracy for Model: 0.9101, Loss for Model: 0.2591

```

Model 8: Added 2 Convolution layers, and 2 Pooling layers, and Fully Connected Layer with a hidden layer (256 Nodes)

```

[35]: model8 = Sequential([
    Input(shape=X_train_4D.shape[1:]),
    Conv2D(32, (3, 3), activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),
    Conv2D(32, (3, 3), activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),
    Flatten(),
    Dense(100, activation='relu'),
    Dense(num_classes, activation='softmax')
])

# Compile the model
model8.compile(loss='categorical_crossentropy', optimizer='adam',
               metrics=['accuracy'])
print(model8.summary())

```

Model: "sequential_9"

Layer (type)	Output Shape	
Param #		
conv2d_2 (Conv2D)	(None, 26, 26, 32)	
↳320		
max_pooling2d_2 (MaxPooling2D)	(None, 13, 13, 32)	
↳ 0		
conv2d_3 (Conv2D)	(None, 11, 11, 32)	
↳9,248		
max_pooling2d_3 (MaxPooling2D)	(None, 5, 5, 32)	
↳ 0		
flatten_8 (Flatten)	(None, 800)	
↳ 0		
dense_24 (Dense)	(None, 100)	
↳80,100		
dense_25 (Dense)	(None, 10)	
↳1,010		

Total params: 90,678 (354.21 KB)

Trainable params: 90,678 (354.21 KB)

Non-trainable params: 0 (0.00 B)

None

```
[36]: # Fit the model
history8 = model8.fit(
    X_train_4D, y_train, validation_data=(X_val_4D, y_val), epochs=100,
    ↳batch_size=2048,
    callbacks=[EarlyStopping(monitor='val_accuracy', patience=5)]
)
```

Epoch 1/100

24/24 9s 296ms/step -

accuracy: 0.3796 - loss: 2.0139 - val_accuracy: 0.6952 - val_loss: 0.8466

Epoch 2/100

24/24 7s 274ms/step -
 accuracy: 0.7109 - loss: 0.7945 - val_accuracy: 0.7453 - val_loss: 0.6768
 Epoch 3/100
 24/24 6s 267ms/step -
 accuracy: 0.7612 - loss: 0.6429 - val_accuracy: 0.7827 - val_loss: 0.5866
 Epoch 4/100
 24/24 6s 264ms/step -
 accuracy: 0.7915 - loss: 0.5646 - val_accuracy: 0.8036 - val_loss: 0.5296
 Epoch 5/100
 24/24 6s 264ms/step -
 accuracy: 0.8122 - loss: 0.5141 - val_accuracy: 0.8172 - val_loss: 0.4984
 Epoch 6/100
 24/24 6s 263ms/step -
 accuracy: 0.8299 - loss: 0.4767 - val_accuracy: 0.8320 - val_loss: 0.4683
 Epoch 7/100
 24/24 6s 263ms/step -
 accuracy: 0.8408 - loss: 0.4517 - val_accuracy: 0.8408 - val_loss: 0.4485
 Epoch 8/100
 24/24 6s 266ms/step -
 accuracy: 0.8470 - loss: 0.4342 - val_accuracy: 0.8453 - val_loss: 0.4330
 Epoch 9/100
 24/24 8s 328ms/step -
 accuracy: 0.8530 - loss: 0.4185 - val_accuracy: 0.8508 - val_loss: 0.4196
 Epoch 10/100
 24/24 9s 358ms/step -
 accuracy: 0.8574 - loss: 0.4045 - val_accuracy: 0.8561 - val_loss: 0.4070
 Epoch 11/100
 24/24 9s 363ms/step -
 accuracy: 0.8627 - loss: 0.3921 - val_accuracy: 0.8587 - val_loss: 0.3976
 Epoch 12/100
 24/24 7s 296ms/step -
 accuracy: 0.8665 - loss: 0.3822 - val_accuracy: 0.8623 - val_loss: 0.3890
 Epoch 13/100
 24/24 6s 257ms/step -
 accuracy: 0.8697 - loss: 0.3731 - val_accuracy: 0.8647 - val_loss: 0.3808
 Epoch 14/100
 24/24 7s 304ms/step -
 accuracy: 0.8727 - loss: 0.3644 - val_accuracy: 0.8681 - val_loss: 0.3724
 Epoch 15/100
 24/24 7s 296ms/step -
 accuracy: 0.8760 - loss: 0.3561 - val_accuracy: 0.8692 - val_loss: 0.3651
 Epoch 16/100
 24/24 6s 262ms/step -
 accuracy: 0.8788 - loss: 0.3486 - val_accuracy: 0.8711 - val_loss: 0.3587
 Epoch 17/100
 24/24 7s 284ms/step -
 accuracy: 0.8803 - loss: 0.3423 - val_accuracy: 0.8724 - val_loss: 0.3531
 Epoch 18/100

24/24 7s 309ms/step -
 accuracy: 0.8815 - loss: 0.3366 - val_accuracy: 0.8748 - val_loss: 0.3473
 Epoch 19/100
 24/24 10s 424ms/step -
 accuracy: 0.8842 - loss: 0.3306 - val_accuracy: 0.8767 - val_loss: 0.3420
 Epoch 20/100
 24/24 8s 327ms/step -
 accuracy: 0.8849 - loss: 0.3251 - val_accuracy: 0.8788 - val_loss: 0.3374
 Epoch 21/100
 24/24 8s 329ms/step -
 accuracy: 0.8866 - loss: 0.3202 - val_accuracy: 0.8799 - val_loss: 0.3337
 Epoch 22/100
 24/24 7s 300ms/step -
 accuracy: 0.8881 - loss: 0.3156 - val_accuracy: 0.8805 - val_loss: 0.3298
 Epoch 23/100
 24/24 7s 283ms/step -
 accuracy: 0.8893 - loss: 0.3109 - val_accuracy: 0.8811 - val_loss: 0.3264
 Epoch 24/100
 24/24 7s 304ms/step -
 accuracy: 0.8910 - loss: 0.3065 - val_accuracy: 0.8817 - val_loss: 0.3234
 Epoch 25/100
 24/24 7s 312ms/step -
 accuracy: 0.8923 - loss: 0.3027 - val_accuracy: 0.8829 - val_loss: 0.3203
 Epoch 26/100
 24/24 7s 299ms/step -
 accuracy: 0.8937 - loss: 0.2989 - val_accuracy: 0.8835 - val_loss: 0.3177
 Epoch 27/100
 24/24 7s 271ms/step -
 accuracy: 0.8951 - loss: 0.2953 - val_accuracy: 0.8852 - val_loss: 0.3149
 Epoch 28/100
 24/24 7s 283ms/step -
 accuracy: 0.8961 - loss: 0.2918 - val_accuracy: 0.8867 - val_loss: 0.3125
 Epoch 29/100
 24/24 7s 306ms/step -
 accuracy: 0.8974 - loss: 0.2884 - val_accuracy: 0.8869 - val_loss: 0.3100
 Epoch 30/100
 24/24 7s 311ms/step -
 accuracy: 0.8989 - loss: 0.2853 - val_accuracy: 0.8878 - val_loss: 0.3078
 Epoch 31/100
 24/24 8s 315ms/step -
 accuracy: 0.9004 - loss: 0.2823 - val_accuracy: 0.8882 - val_loss: 0.3058
 Epoch 32/100
 24/24 7s 293ms/step -
 accuracy: 0.9010 - loss: 0.2793 - val_accuracy: 0.8895 - val_loss: 0.3034
 Epoch 33/100
 24/24 7s 277ms/step -
 accuracy: 0.9019 - loss: 0.2765 - val_accuracy: 0.8898 - val_loss: 0.3014
 Epoch 34/100

24/24 7s 275ms/step -
 accuracy: 0.9035 - loss: 0.2736 - val_accuracy: 0.8905 - val_loss: 0.2995
 Epoch 35/100
 24/24 7s 277ms/step -
 accuracy: 0.9041 - loss: 0.2707 - val_accuracy: 0.8912 - val_loss: 0.2978
 Epoch 36/100
 24/24 7s 304ms/step -
 accuracy: 0.9053 - loss: 0.2681 - val_accuracy: 0.8915 - val_loss: 0.2959
 Epoch 37/100
 24/24 7s 291ms/step -
 accuracy: 0.9060 - loss: 0.2654 - val_accuracy: 0.8920 - val_loss: 0.2938
 Epoch 38/100
 24/24 7s 275ms/step -
 accuracy: 0.9067 - loss: 0.2626 - val_accuracy: 0.8934 - val_loss: 0.2925
 Epoch 39/100
 24/24 7s 293ms/step -
 accuracy: 0.9075 - loss: 0.2603 - val_accuracy: 0.8933 - val_loss: 0.2909
 Epoch 40/100
 24/24 7s 305ms/step -
 accuracy: 0.9086 - loss: 0.2578 - val_accuracy: 0.8932 - val_loss: 0.2902
 Epoch 41/100
 24/24 7s 292ms/step -
 accuracy: 0.9092 - loss: 0.2559 - val_accuracy: 0.8935 - val_loss: 0.2885
 Epoch 42/100
 24/24 7s 293ms/step -
 accuracy: 0.9098 - loss: 0.2535 - val_accuracy: 0.8947 - val_loss: 0.2869
 Epoch 43/100
 24/24 8s 326ms/step -
 accuracy: 0.9104 - loss: 0.2512 - val_accuracy: 0.8950 - val_loss: 0.2858
 Epoch 44/100
 24/24 7s 273ms/step -
 accuracy: 0.9111 - loss: 0.2490 - val_accuracy: 0.8956 - val_loss: 0.2846
 Epoch 45/100
 24/24 7s 288ms/step -
 accuracy: 0.9121 - loss: 0.2470 - val_accuracy: 0.8964 - val_loss: 0.2834
 Epoch 46/100
 24/24 7s 277ms/step -
 accuracy: 0.9130 - loss: 0.2447 - val_accuracy: 0.8968 - val_loss: 0.2821
 Epoch 47/100
 24/24 8s 327ms/step -
 accuracy: 0.9140 - loss: 0.2425 - val_accuracy: 0.8975 - val_loss: 0.2813
 Epoch 48/100
 24/24 8s 340ms/step -
 accuracy: 0.9147 - loss: 0.2406 - val_accuracy: 0.8983 - val_loss: 0.2802
 Epoch 49/100
 24/24 8s 315ms/step -
 accuracy: 0.9153 - loss: 0.2386 - val_accuracy: 0.8987 - val_loss: 0.2794
 Epoch 50/100

24/24 7s 276ms/step -
 accuracy: 0.9160 - loss: 0.2367 - val_accuracy: 0.8986 - val_loss: 0.2786
 Epoch 51/100
 24/24 7s 271ms/step -
 accuracy: 0.9163 - loss: 0.2349 - val_accuracy: 0.8992 - val_loss: 0.2778
 Epoch 52/100
 24/24 7s 269ms/step -
 accuracy: 0.9169 - loss: 0.2330 - val_accuracy: 0.8994 - val_loss: 0.2769
 Epoch 53/100
 24/24 7s 272ms/step -
 accuracy: 0.9176 - loss: 0.2312 - val_accuracy: 0.8992 - val_loss: 0.2766
 Epoch 54/100
 24/24 7s 294ms/step -
 accuracy: 0.9184 - loss: 0.2296 - val_accuracy: 0.9003 - val_loss: 0.2760
 Epoch 55/100
 24/24 7s 270ms/step -
 accuracy: 0.9189 - loss: 0.2280 - val_accuracy: 0.9008 - val_loss: 0.2753
 Epoch 56/100
 24/24 6s 269ms/step -
 accuracy: 0.9194 - loss: 0.2264 - val_accuracy: 0.9014 - val_loss: 0.2745
 Epoch 57/100
 24/24 7s 281ms/step -
 accuracy: 0.9202 - loss: 0.2247 - val_accuracy: 0.9017 - val_loss: 0.2744
 Epoch 58/100
 24/24 7s 270ms/step -
 accuracy: 0.9206 - loss: 0.2233 - val_accuracy: 0.9012 - val_loss: 0.2732
 Epoch 59/100
 24/24 7s 283ms/step -
 accuracy: 0.9218 - loss: 0.2213 - val_accuracy: 0.9008 - val_loss: 0.2730
 Epoch 60/100
 24/24 7s 286ms/step -
 accuracy: 0.9224 - loss: 0.2198 - val_accuracy: 0.9012 - val_loss: 0.2725
 Epoch 61/100
 24/24 6s 268ms/step -
 accuracy: 0.9224 - loss: 0.2182 - val_accuracy: 0.9018 - val_loss: 0.2724
 Epoch 62/100
 24/24 7s 285ms/step -
 accuracy: 0.9223 - loss: 0.2173 - val_accuracy: 0.9015 - val_loss: 0.2730
 Epoch 63/100
 24/24 7s 279ms/step -
 accuracy: 0.9217 - loss: 0.2168 - val_accuracy: 0.9003 - val_loss: 0.2759
 Epoch 64/100
 24/24 7s 291ms/step -
 accuracy: 0.9205 - loss: 0.2181 - val_accuracy: 0.8957 - val_loss: 0.2846
 Epoch 65/100
 24/24 7s 283ms/step -
 accuracy: 0.9186 - loss: 0.2218 - val_accuracy: 0.9029 - val_loss: 0.2686
 Epoch 66/100

```

24/24          7s 286ms/step -
accuracy: 0.9242 - loss: 0.2120 - val_accuracy: 0.9016 - val_loss: 0.2673
Epoch 67/100
24/24          7s 280ms/step -
accuracy: 0.9259 - loss: 0.2073 - val_accuracy: 0.9017 - val_loss: 0.2668
Epoch 68/100
24/24          7s 287ms/step -
accuracy: 0.9263 - loss: 0.2061 - val_accuracy: 0.9019 - val_loss: 0.2666
Epoch 69/100
24/24          7s 301ms/step -
accuracy: 0.9265 - loss: 0.2046 - val_accuracy: 0.9022 - val_loss: 0.2657
Epoch 70/100
24/24          7s 277ms/step -
accuracy: 0.9273 - loss: 0.2027 - val_accuracy: 0.9020 - val_loss: 0.2659

```

```

[37]: # Evaluation of the model on the validation set
scores8 = model8.evaluate(X_val_4D, y_val)

print("\n")
print(f"Accuracy for Model 8: {round(scores8[1], 4)}, Loss for Model 8: {round(scores8[0], 4)}")

```

```

375/375        2s 6ms/step -
accuracy: 0.9026 - loss: 0.2604

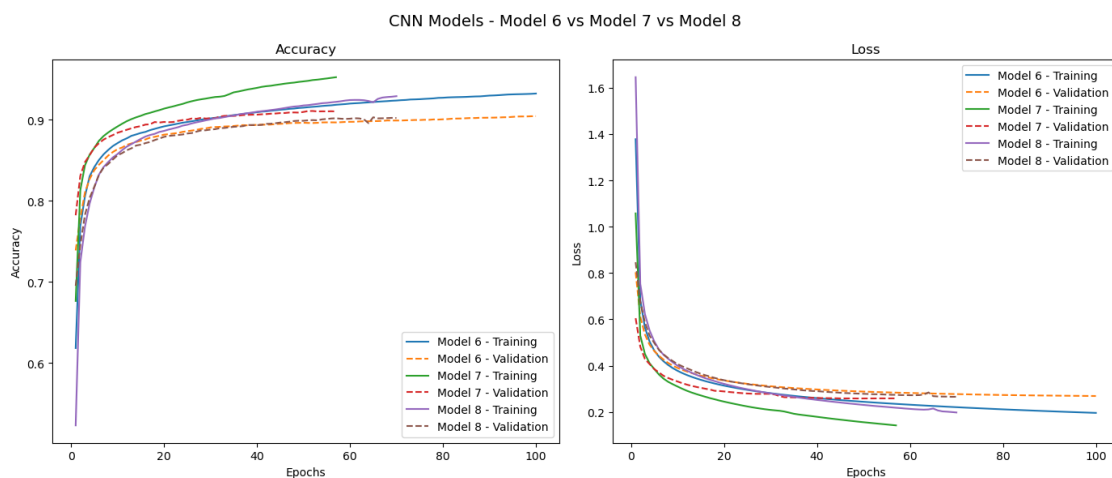
```

Accuracy for Model 8: 0.902, Loss for Model 8: 0.2657

```

[99]: plot_model_history([history6, history7, history8], ['Model 6', 'Model 7', 'Model 8'], 'CNN Models - Model 6 vs Model 7 vs Model 8')

```



Model 6 has no middle layer, while model 7 has a middle layer with 100 nodes. Model 8 on the

other hand has two convolution layers with a layer of 100 nodes. It is very clear that model 7 is by far the best on validation sets

0.9 6. Try to use a pre-trained network to improve accuracy.

```
[38]: from keras.applications.resnet50 import ResNet50, preprocess_input
      from keras.utils import load_img, img_to_array
      from keras.utils import to_categorical
      from keras.models import Sequential
      from keras.layers import Dense, Flatten, Input
      from skimage.color import gray2rgb
      from skimage.transform import resize

[60]: # Preprocess the data and convert grayscale images to RGB
      X_train_resized = np.array([preprocess_input(gray2rgb(img)) for img in X_train])
      X_test_resized = np.array([preprocess_input(gray2rgb(img)) for img in X_test])

      # Resize the images to match the input shape expected by ResNet50
      X_train_resized = np.array([resize(img, (32, 32)) for img in X_train])
      X_val_resized = np.array([resize(img, (32, 32)) for img in X_val])

      # Load pre-trained ResNet50 model without the top layer (include_top=False)
      resnet_model = ResNet50(weights='imagenet', include_top=False, input_shape=(32, 32, 3))

      # Freeze the layers in the pre-trained model
      for layer in resnet_model.layers:
          layer.trainable = False

      # Create a new model and add the pre-trained ResNet50 base
      pretrained_model = Sequential()
      pretrained_model.add(resnet_model)

      # Add additional layers on top of the ResNet50 base
      pretrained_model.add(Flatten())
      pretrained_model.add(Dense(100, activation='relu'))
      pretrained_model.add(Dense(10, activation='softmax')) # FashionMNIST has 10 classes

      # Compile the model
      pretrained_model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
      print(pretrained_model.summary())
```

Model: "sequential_17"

Layer (type)	Output Shape	
↳Param #		
resnet50 (Functional)	?	
↳23,587,712		
flatten_11 (Flatten)	?	0
↳(unbuilt)		
dense_30 (Dense)	?	0
↳(unbuilt)		
dense_31 (Dense)	?	0
↳(unbuilt)		

Total params: 23,587,712 (89.98 MB)

Trainable params: 0 (0.00 B)

Non-trainable params: 23,587,712 (89.98 MB)

None

```
[44]: # Train the model
history_pre = pretrained_model.fit(X_train_resized, y_train,
↳validation_data=(X_val_resized, y_val), epochs=20, batch_size=4096,
↳callbacks=[EarlyStopping(monitor='val_accuracy', patience=5)])
```

Epoch 1/20

12/12 107s 8s/step -

accuracy: 0.4011 - loss: 2.8412 - val_accuracy: 0.7360 - val_loss: 0.8053

Epoch 2/20

12/12 79s 7s/step -

accuracy: 0.7564 - loss: 0.7320 - val_accuracy: 0.7747 - val_loss: 0.6158

Epoch 3/20

12/12 80s 7s/step -

accuracy: 0.7915 - loss: 0.5784 - val_accuracy: 0.8033 - val_loss: 0.5363

Epoch 4/20

12/12 80s 7s/step -

accuracy: 0.8137 - loss: 0.5093 - val_accuracy: 0.8183 - val_loss: 0.4959

Epoch 5/20

12/12 81s 7s/step -


accuracy: 0.8319 - loss: 0.4660 - val_accuracy: 0.8270 - val_loss: 0.4682

Epoch 6/20
12/12 81s 7s/step -
accuracy: 0.8424 - loss: 0.4333 - val_accuracy: 0.8375 - val_loss: 0.4480
Epoch 7/20
12/12 82s 7s/step -
accuracy: 0.8525 - loss: 0.4067 - val_accuracy: 0.8421 - val_loss: 0.4336
Epoch 8/20
12/12 89s 7s/step -
accuracy: 0.8605 - loss: 0.3853 - val_accuracy: 0.8451 - val_loss: 0.4227
Epoch 9/20
12/12 80s 7s/step -
accuracy: 0.8677 - loss: 0.3667 - val_accuracy: 0.8491 - val_loss: 0.4142
Epoch 10/20
12/12 88s 7s/step -
accuracy: 0.8737 - loss: 0.3509 - val_accuracy: 0.8515 - val_loss: 0.4075
Epoch 11/20
12/12 83s 7s/step -
accuracy: 0.8793 - loss: 0.3368 - val_accuracy: 0.8533 - val_loss: 0.4018
Epoch 12/20
12/12 80s 7s/step -
accuracy: 0.8836 - loss: 0.3240 - val_accuracy: 0.8546 - val_loss: 0.3974
Epoch 13/20
12/12 86s 7s/step -
accuracy: 0.8880 - loss: 0.3123 - val_accuracy: 0.8572 - val_loss: 0.3929
Epoch 14/20
12/12 77s 6s/step -
accuracy: 0.8928 - loss: 0.3016 - val_accuracy: 0.8593 - val_loss: 0.3891
Epoch 15/20
12/12 78s 7s/step -
accuracy: 0.8969 - loss: 0.2917 - val_accuracy: 0.8599 - val_loss: 0.3857
Epoch 16/20
12/12 81s 7s/step -
accuracy: 0.9009 - loss: 0.2825 - val_accuracy: 0.8608 - val_loss: 0.3837
Epoch 17/20
12/12 77s 7s/step -
accuracy: 0.9037 - loss: 0.2743 - val_accuracy: 0.8618 - val_loss: 0.3814
Epoch 18/20
12/12 78s 7s/step -
accuracy: 0.9071 - loss: 0.2663 - val_accuracy: 0.8620 - val_loss: 0.3805
Epoch 19/20
12/12 79s 7s/step -
accuracy: 0.9098 - loss: 0.2590 - val_accuracy: 0.8626 - val_loss: 0.3798
Epoch 20/20
12/12 80s 7s/step -
accuracy: 0.9128 - loss: 0.2520 - val_accuracy: 0.8617 - val_loss: 0.3797

```
[46]: # Evaluation of the model on the validation set
scores_pre = pretrained_model.evaluate(X_val_resized, y_val)
print(f"Accuracy for ResNet50 : {round(scores_pre[1], 4)}")
```

```
375/375          31s 84ms/step -
accuracy: 0.8643 - loss: 0.3717
Accuracy for ResNet50 : 0.8617
```

```
[100]: plot_model_history([history7,history_pre], ['Model 7','ResNet50'],'Comparison -_
↳Model 7 vs Pre Trained Model')
```



Using ResNet50 as a pretrained model did not improve accuracy. One reason could be as we tried changin greyscale images to rgb which ResNet50 requires.

0.10 7. Select a final model and evaluate it on the test set. How does the test error compare to the validation error?

My best model was **Model 7: Adding Convolution, Pooling, and Fully Connected Layer with a hidden layer (100 Nodes)**

```
[70]: print(f"Validation Accuracy: {round(scores7[1], 4)}, Validation Loss:_)
↳{round(scores7[0], 4)}")
print("\n")
```

```
Accuracy for Model: 0.9101, Loss for Model: 0.2591
```

```
[71]: scores7_test = model7.evaluate(X_test_4D, y_test)
```

313/313 1s 3ms/step -
accuracy: 0.9104 - loss: 0.2754

```
[73]: print(f"Test Accuracy: {round(scores7_test[1], 4)}, Test Loss: ␣  
      ↪ {round(scores7_test[0], 4)}")  
      print("\n")
```

Test Accuracy: 0.9084, Test Loss: 0.2681

```
[106]: data = {  
        "Model 7": ["Validation", "Test"],  
        "Accuracy": [round(scores7[1], 4), round(scores7_test[1], 4)],  
        "Loss": [round(scores7[0], 4), round(scores7_test[0], 4)]  
      }  
  
      # Create a dataframe from the dictionary  
      df = pd.DataFrame(data)  
  
      df
```

```
[106]:      Model 7  Accuracy  Loss  
0  Validation    0.9101  0.2591  
1         Test    0.9084  0.2681
```

As you can see the accuracy for test dataset for our best model is pretty close to the validation set as well as the loss score. For reminder, find the details of the model below

```
[108]: print(model7.summary())
```

Model: "sequential_8"

Layer (type)	Output Shape	
↪ Param #		
conv2d_1 (Conv2D)	(None, 26, 26, 32)	␣
↪ 320		
max_pooling2d_1 (MaxPooling2D)	(None, 13, 13, 32)	␣
↪ 0		
flatten_7 (Flatten)	(None, 5408)	␣
↪ 0		

```
dense_22 (Dense)                (None, 100)
↳540,900                        ↳
dense_23 (Dense)                (None, 10)
↳1,010                          ↳
```

Total params: 1,626,692 (6.21 MB)

Trainable params: 542,230 (2.07 MB)

Non-trainable params: 0 (0.00 B)

Optimizer params: 1,084,462 (4.14 MB)

None

[]: